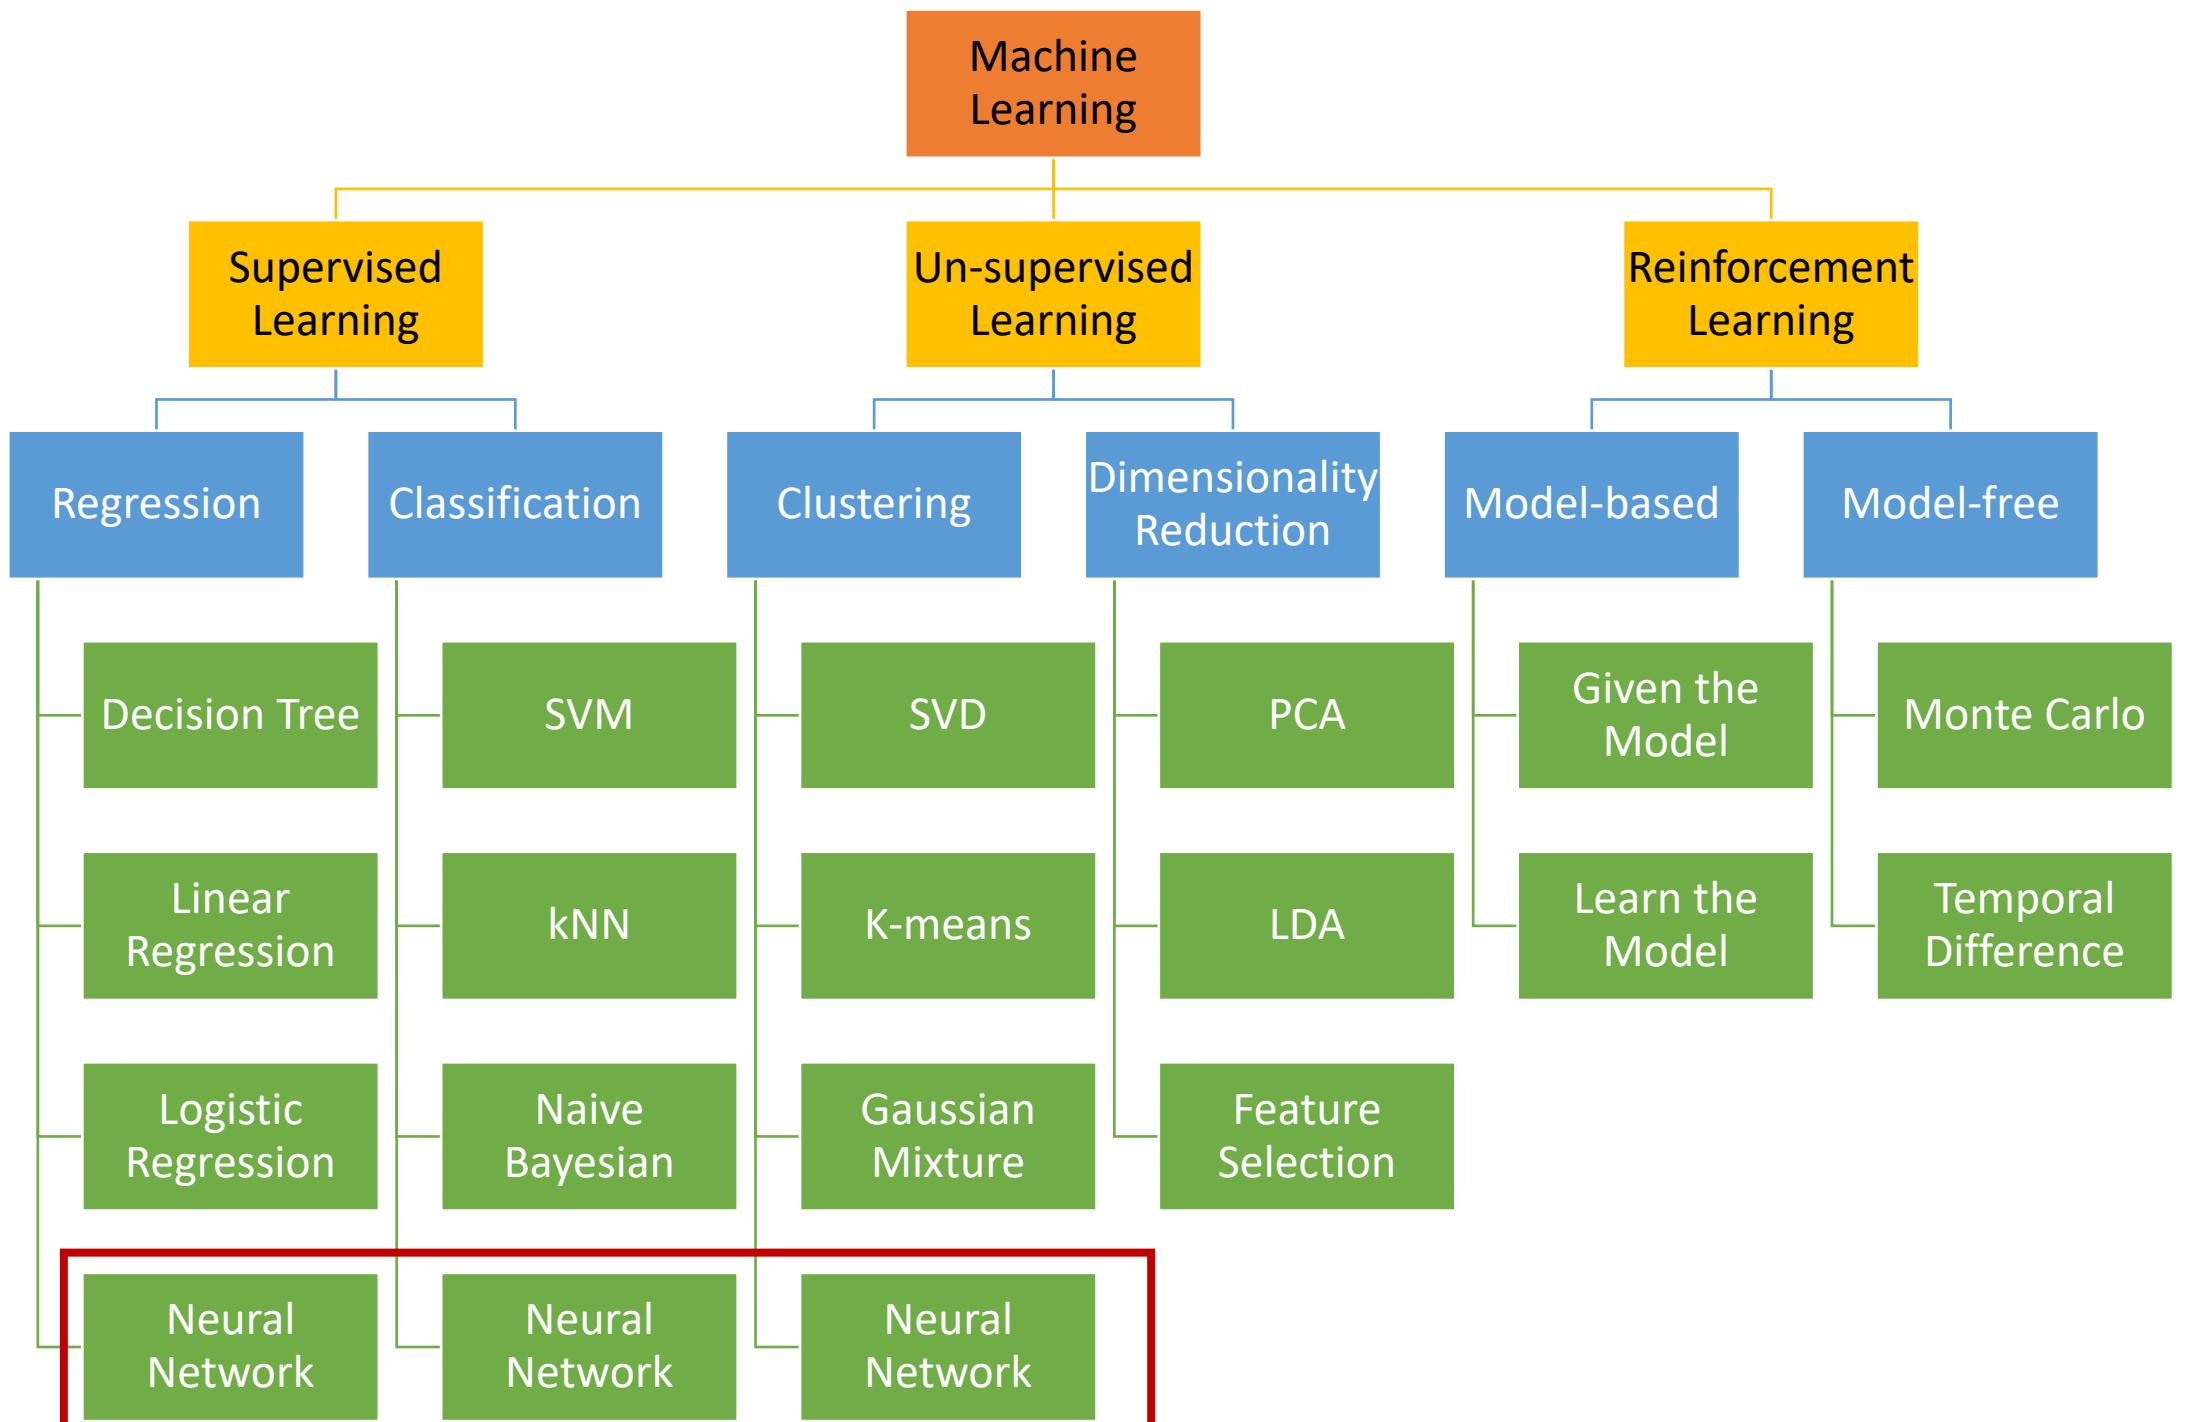


Summer Bootcamp 2021

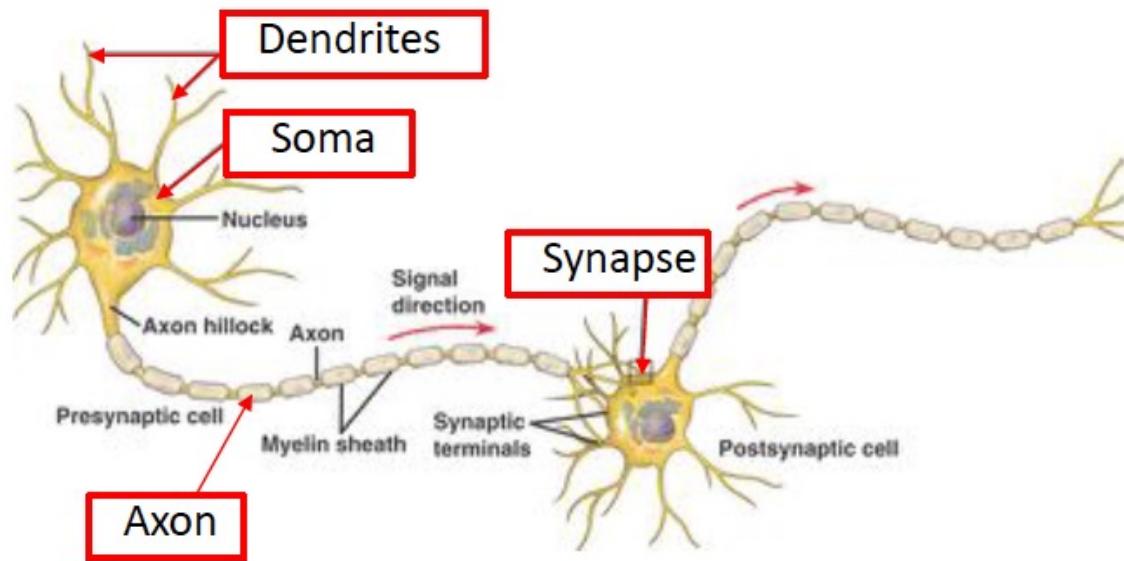
Applied Machine Learning Convolutional Neural Networks

Ngan Le
thile@uark.edu



Perceptron

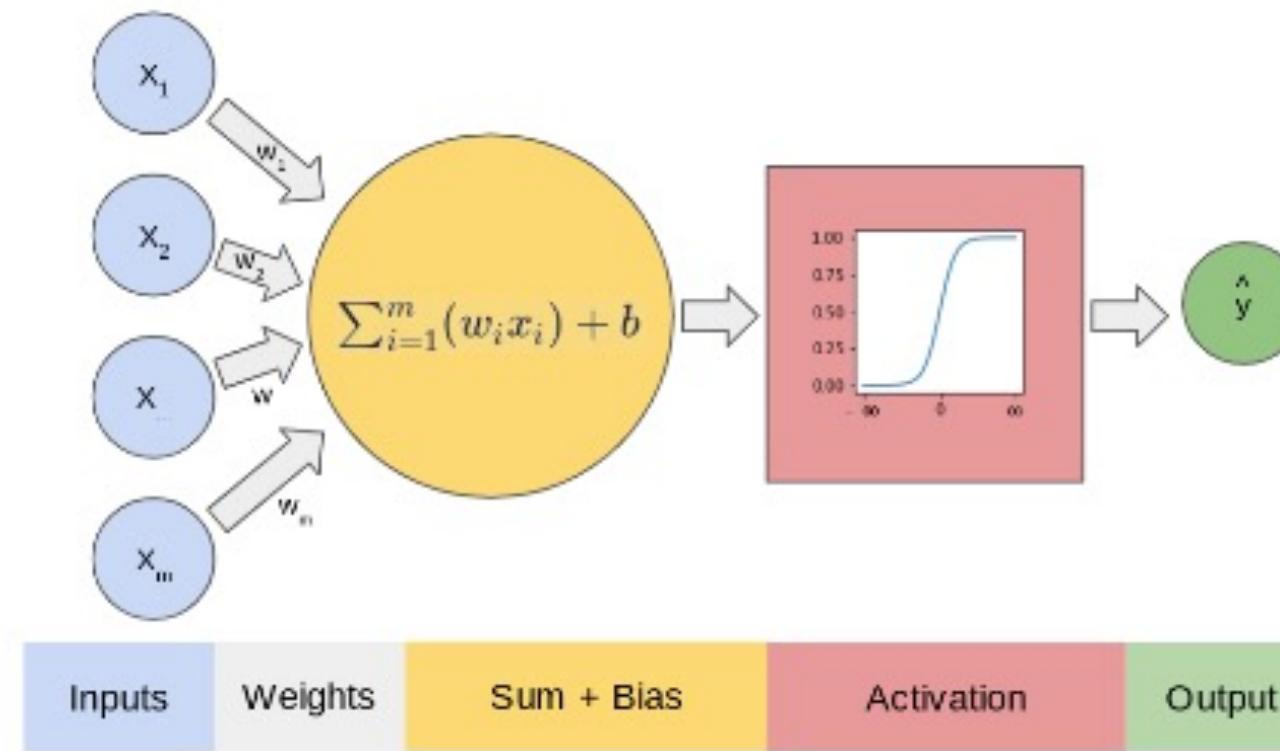
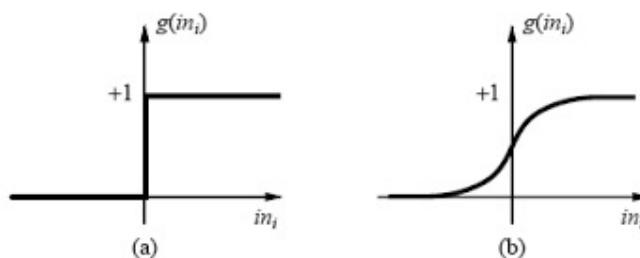
- A Neuron:



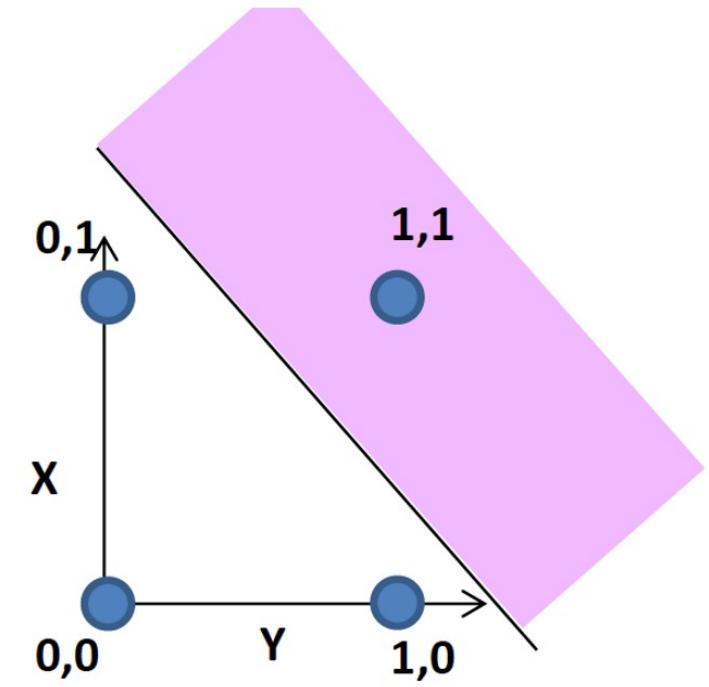
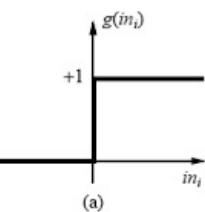
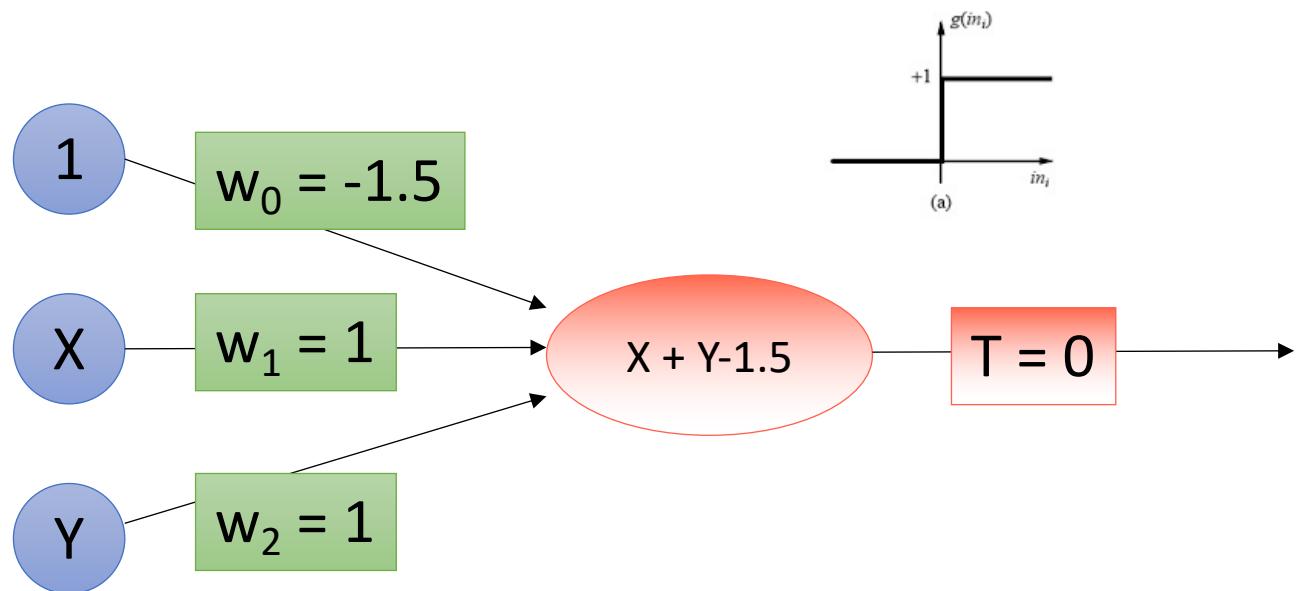
If two neurons consistently fire simultaneously, synaptic connection is increased, (if firing at different time, strength is reduced)

Cells that fire together, wire together

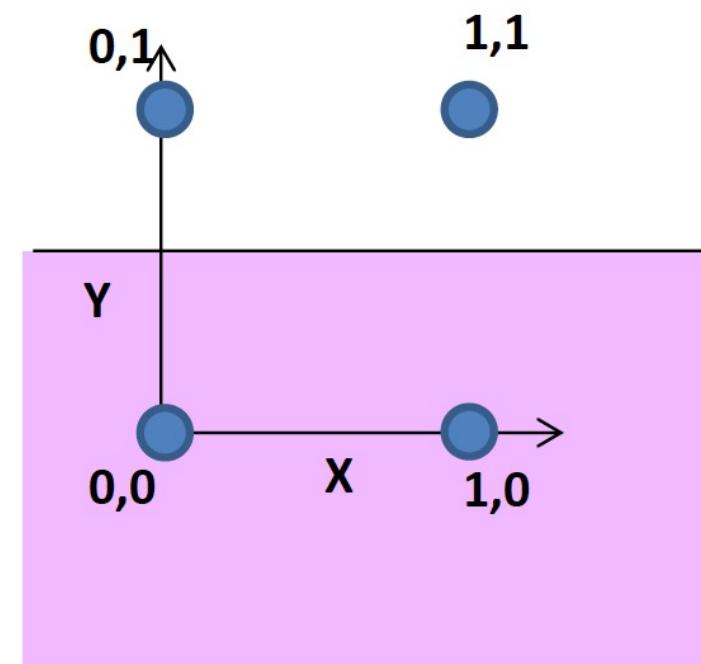
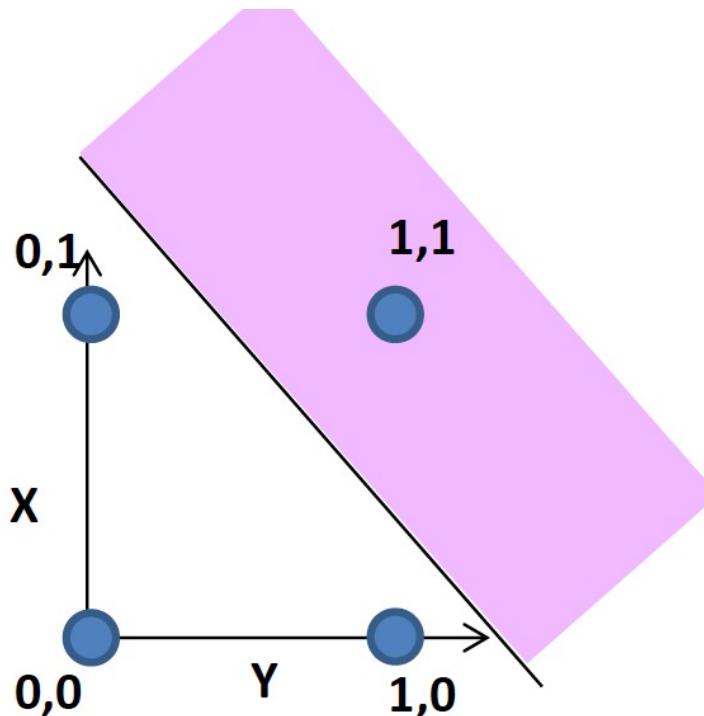
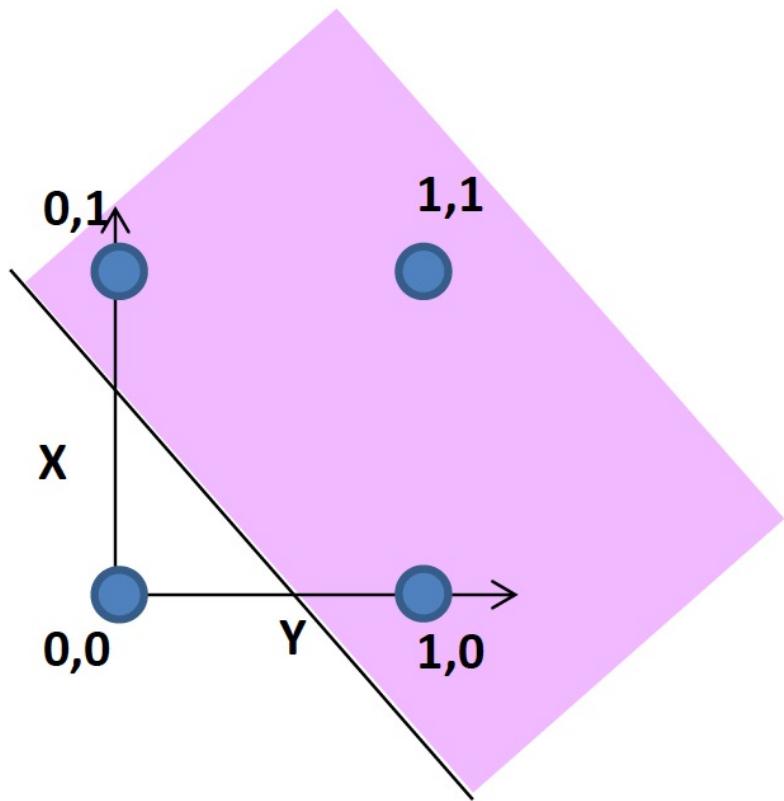
Perceptron



Perceptron

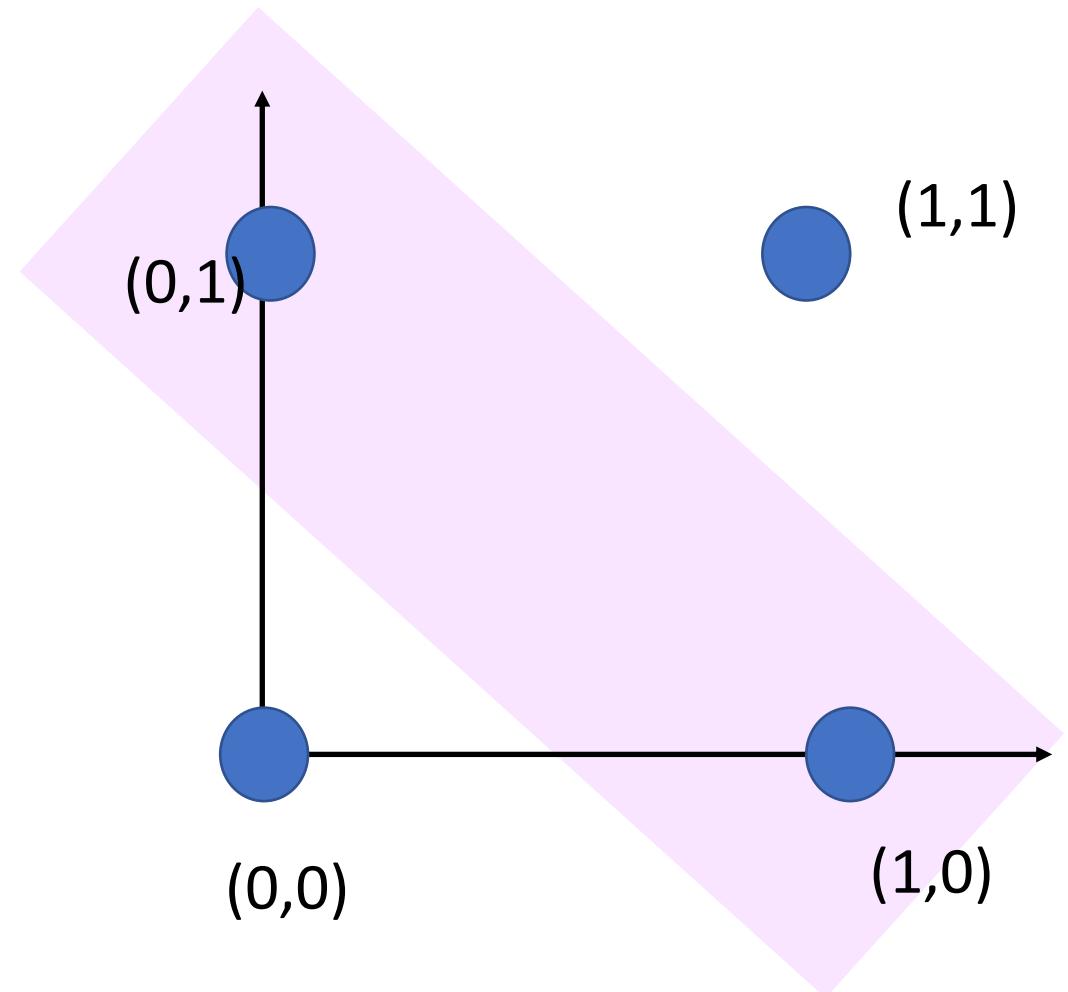


Perceptron

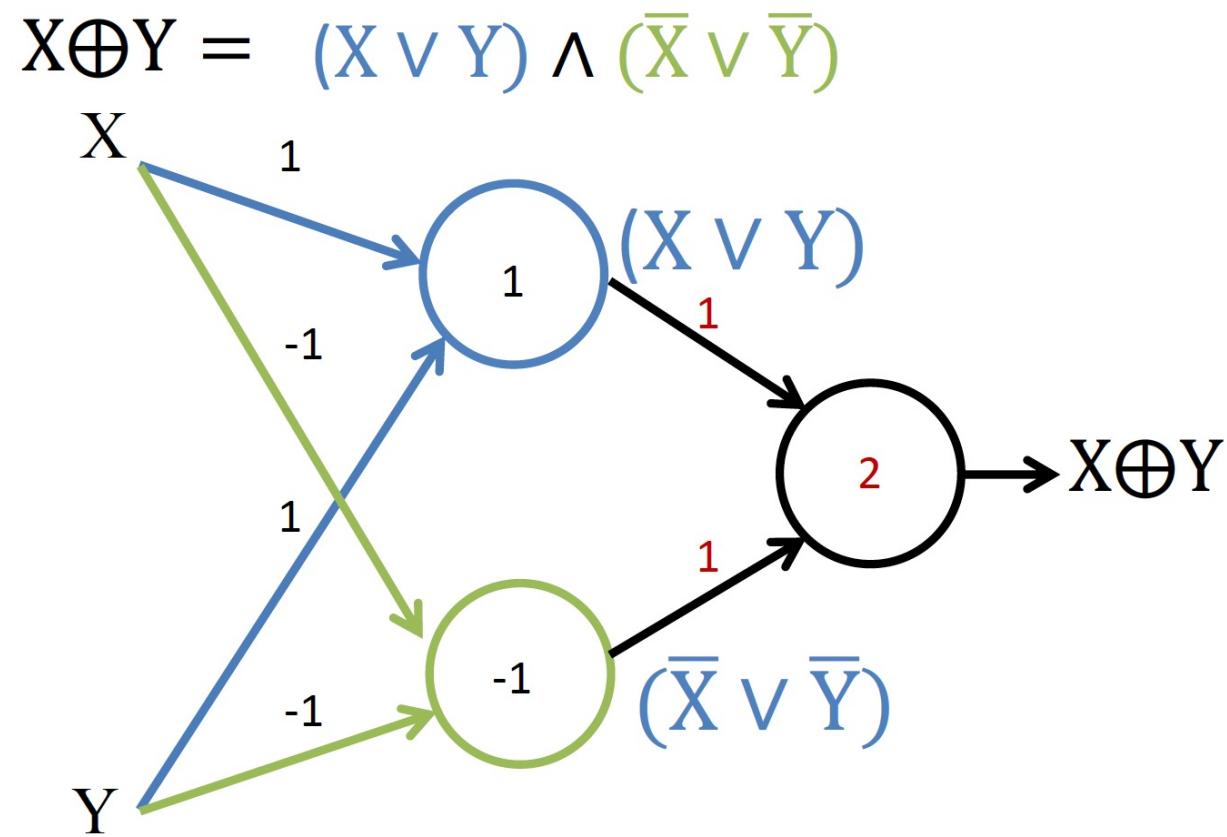


Neuron - Perceptron

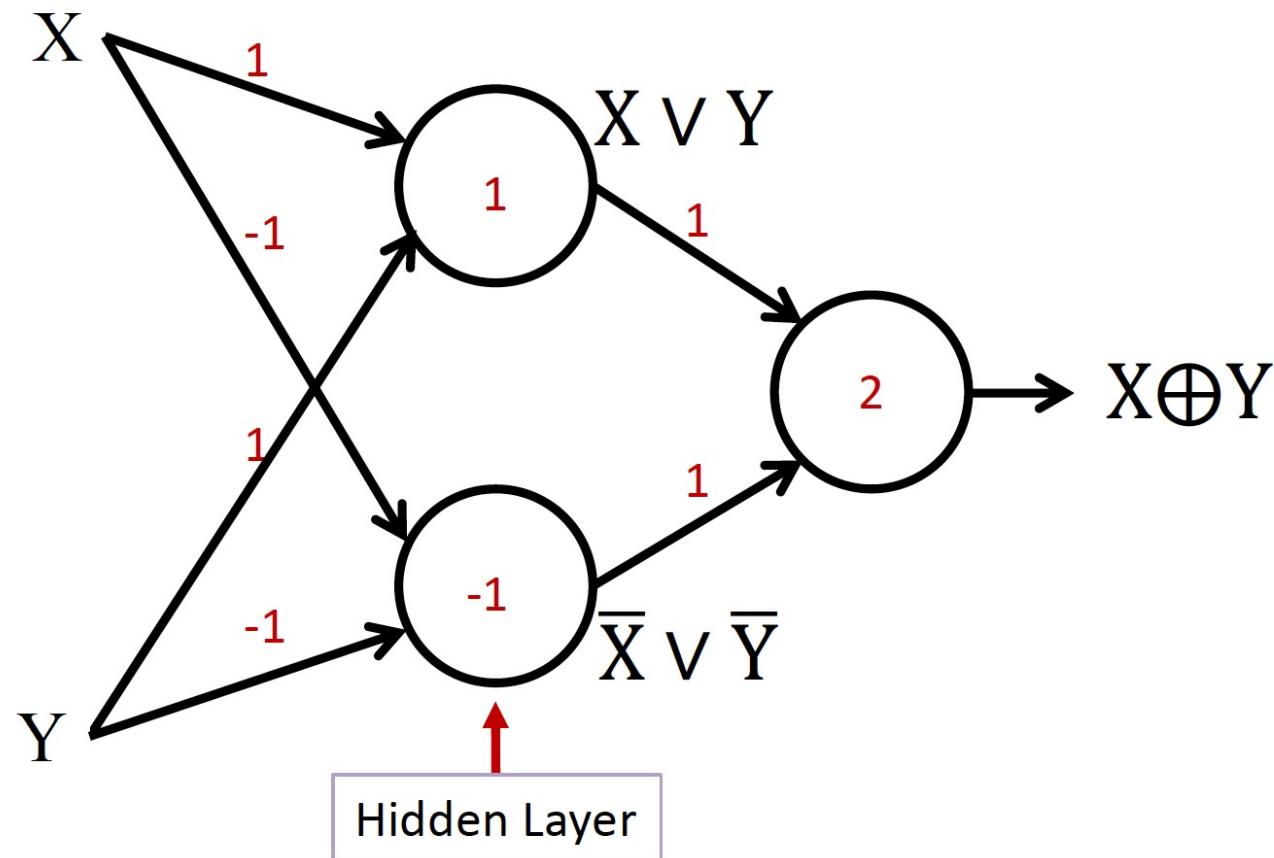
x_1	x_2	$x_1 \text{ XOR } x_2$
0	0	0
0	1	1
1	0	1
1	1	0



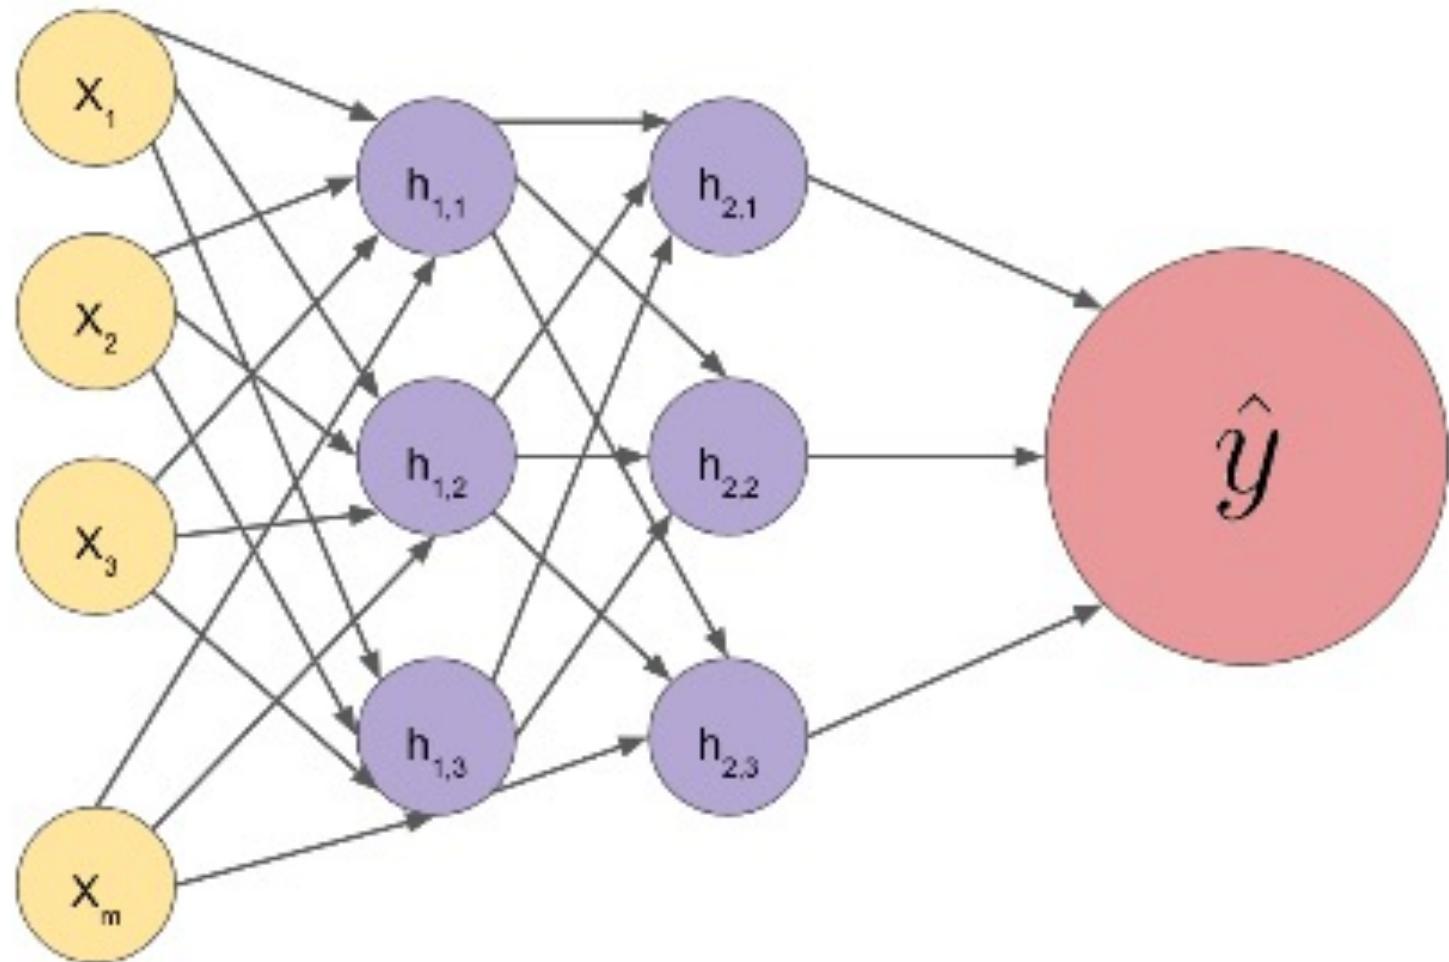
Multilayer Perceptron (MLP)



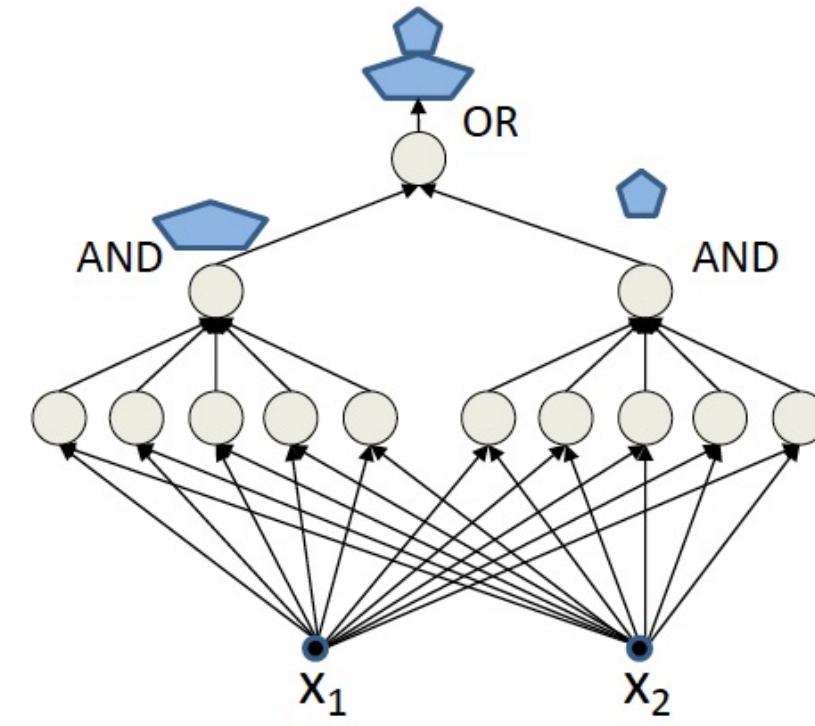
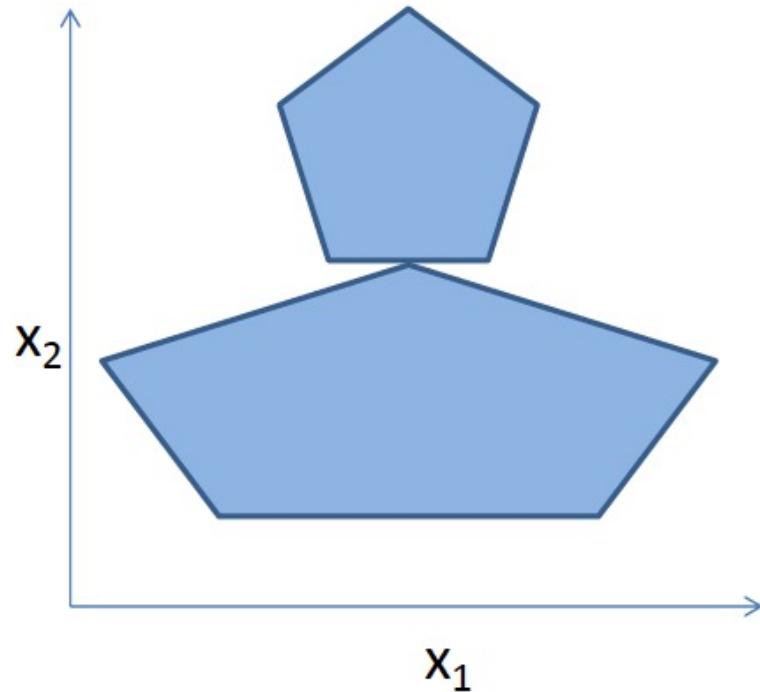
Multilayer Perceptron (MLP)



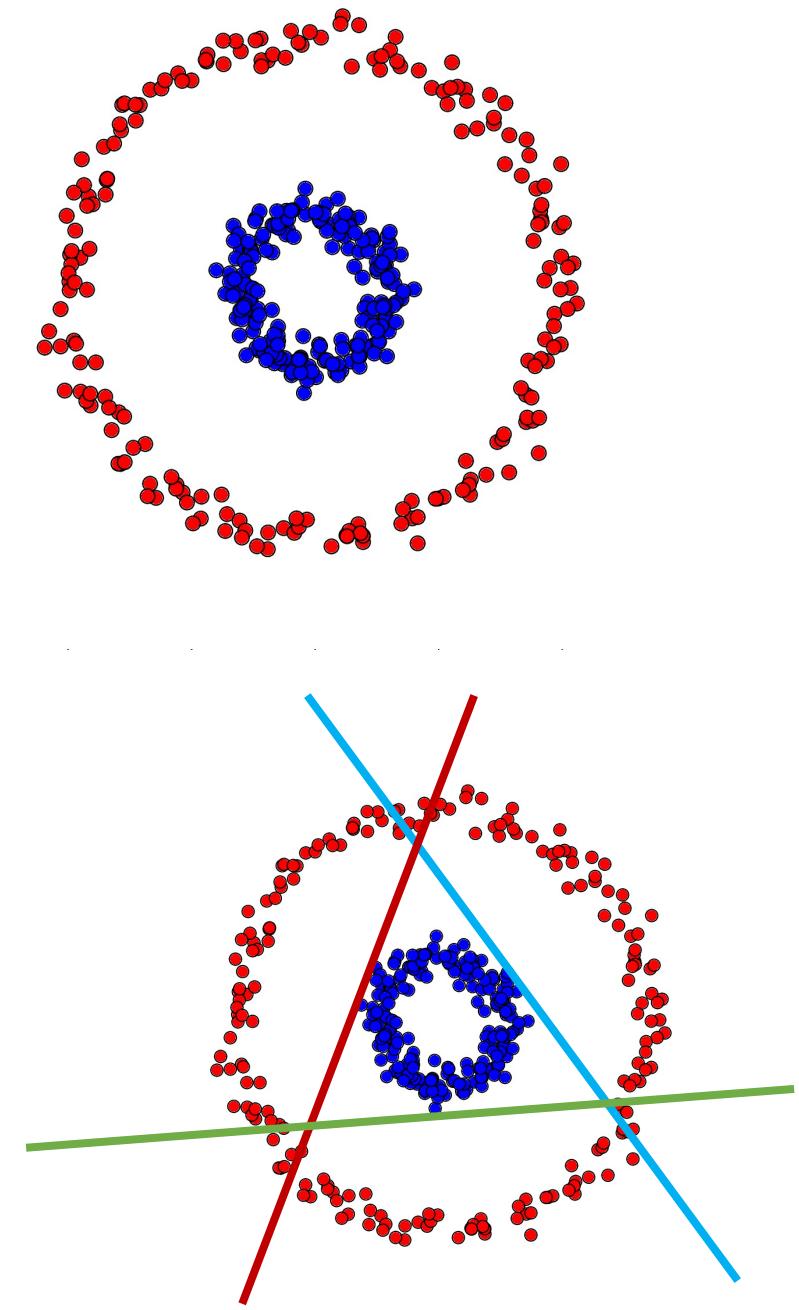
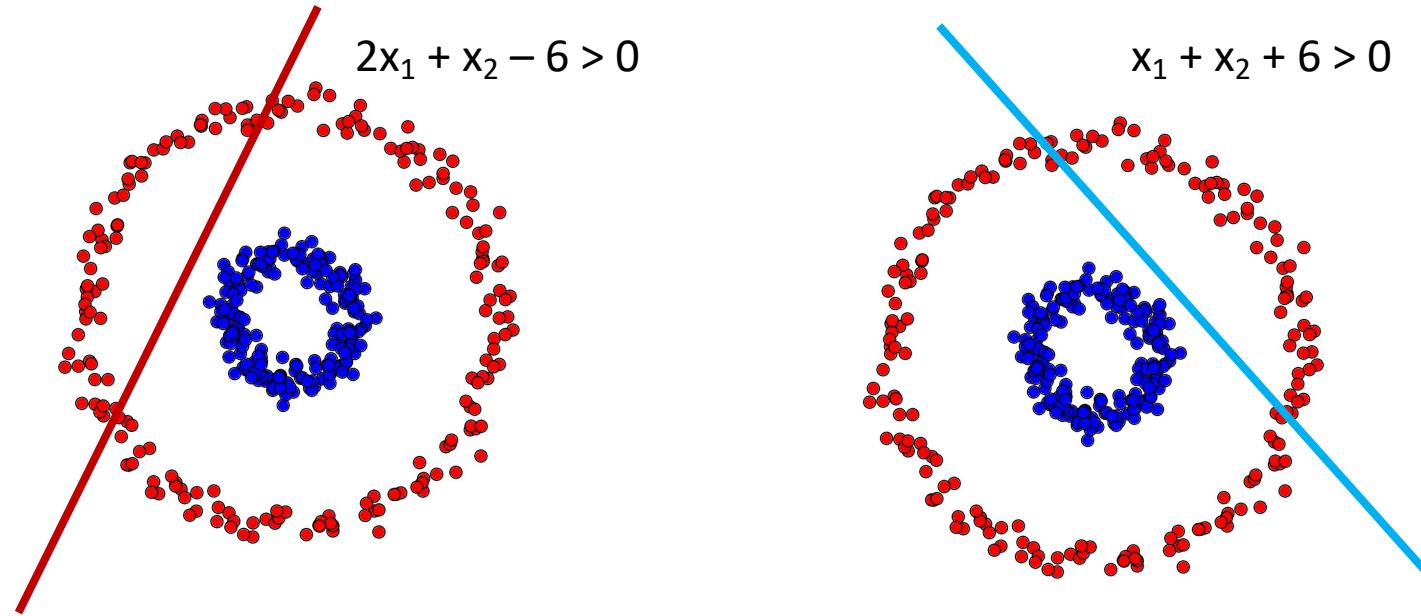
Neuron with Hidden Layers



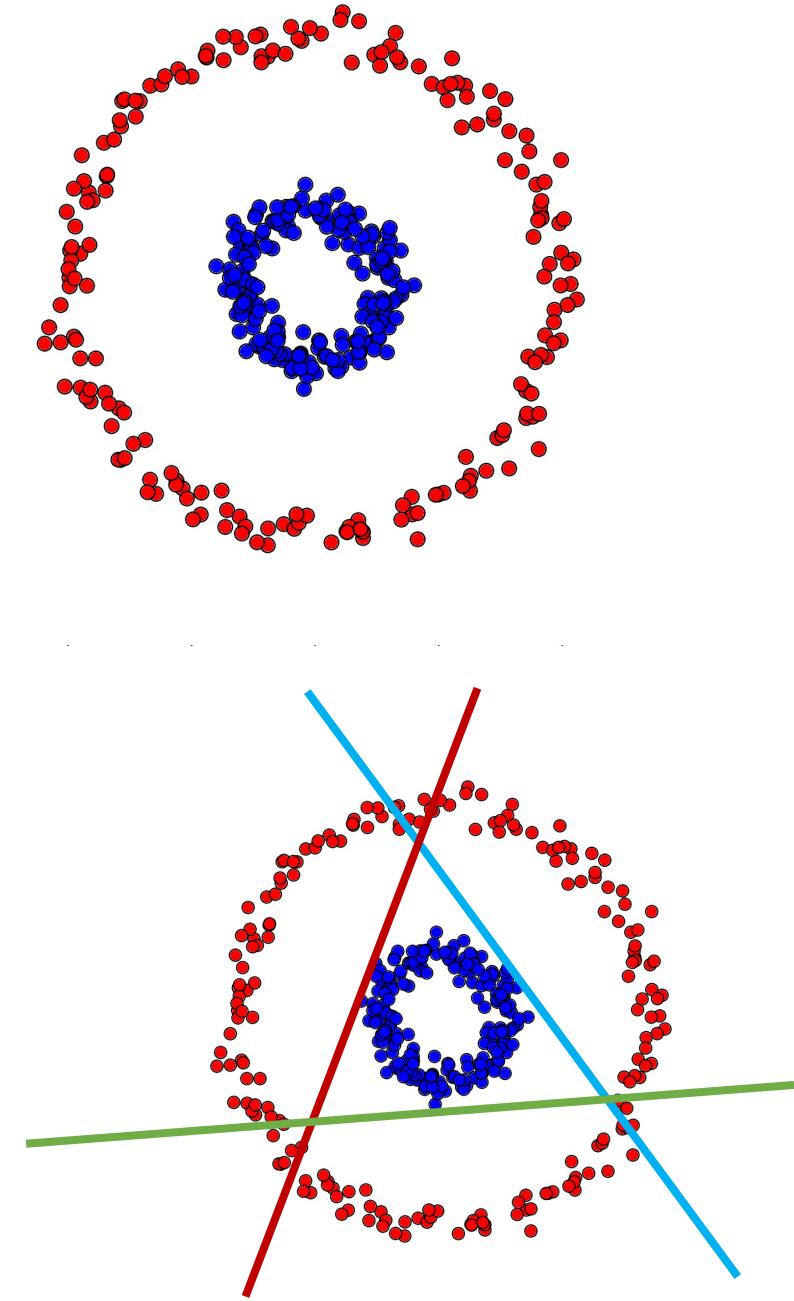
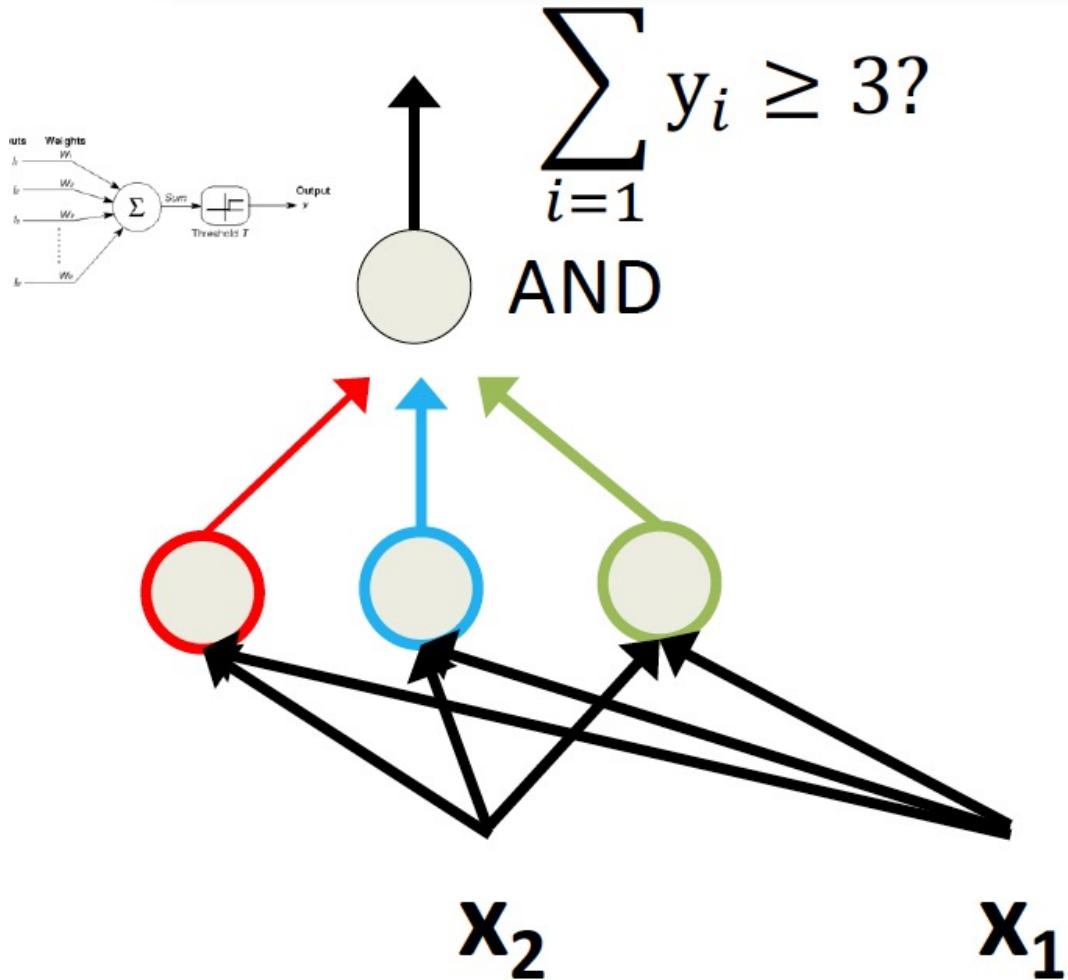
Neuron with Hidden Layers



MLP for Classification

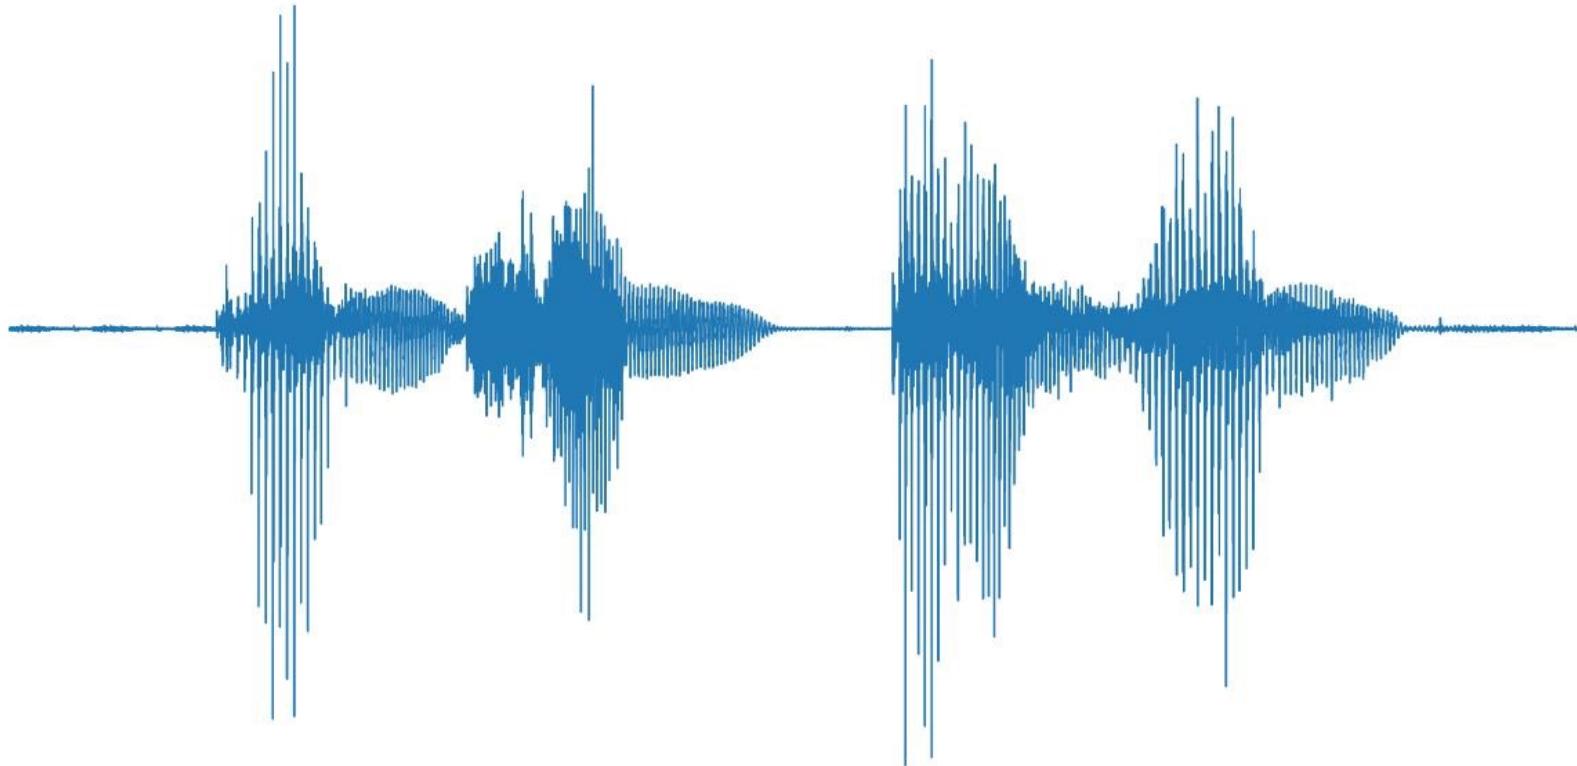


MLP for Classification



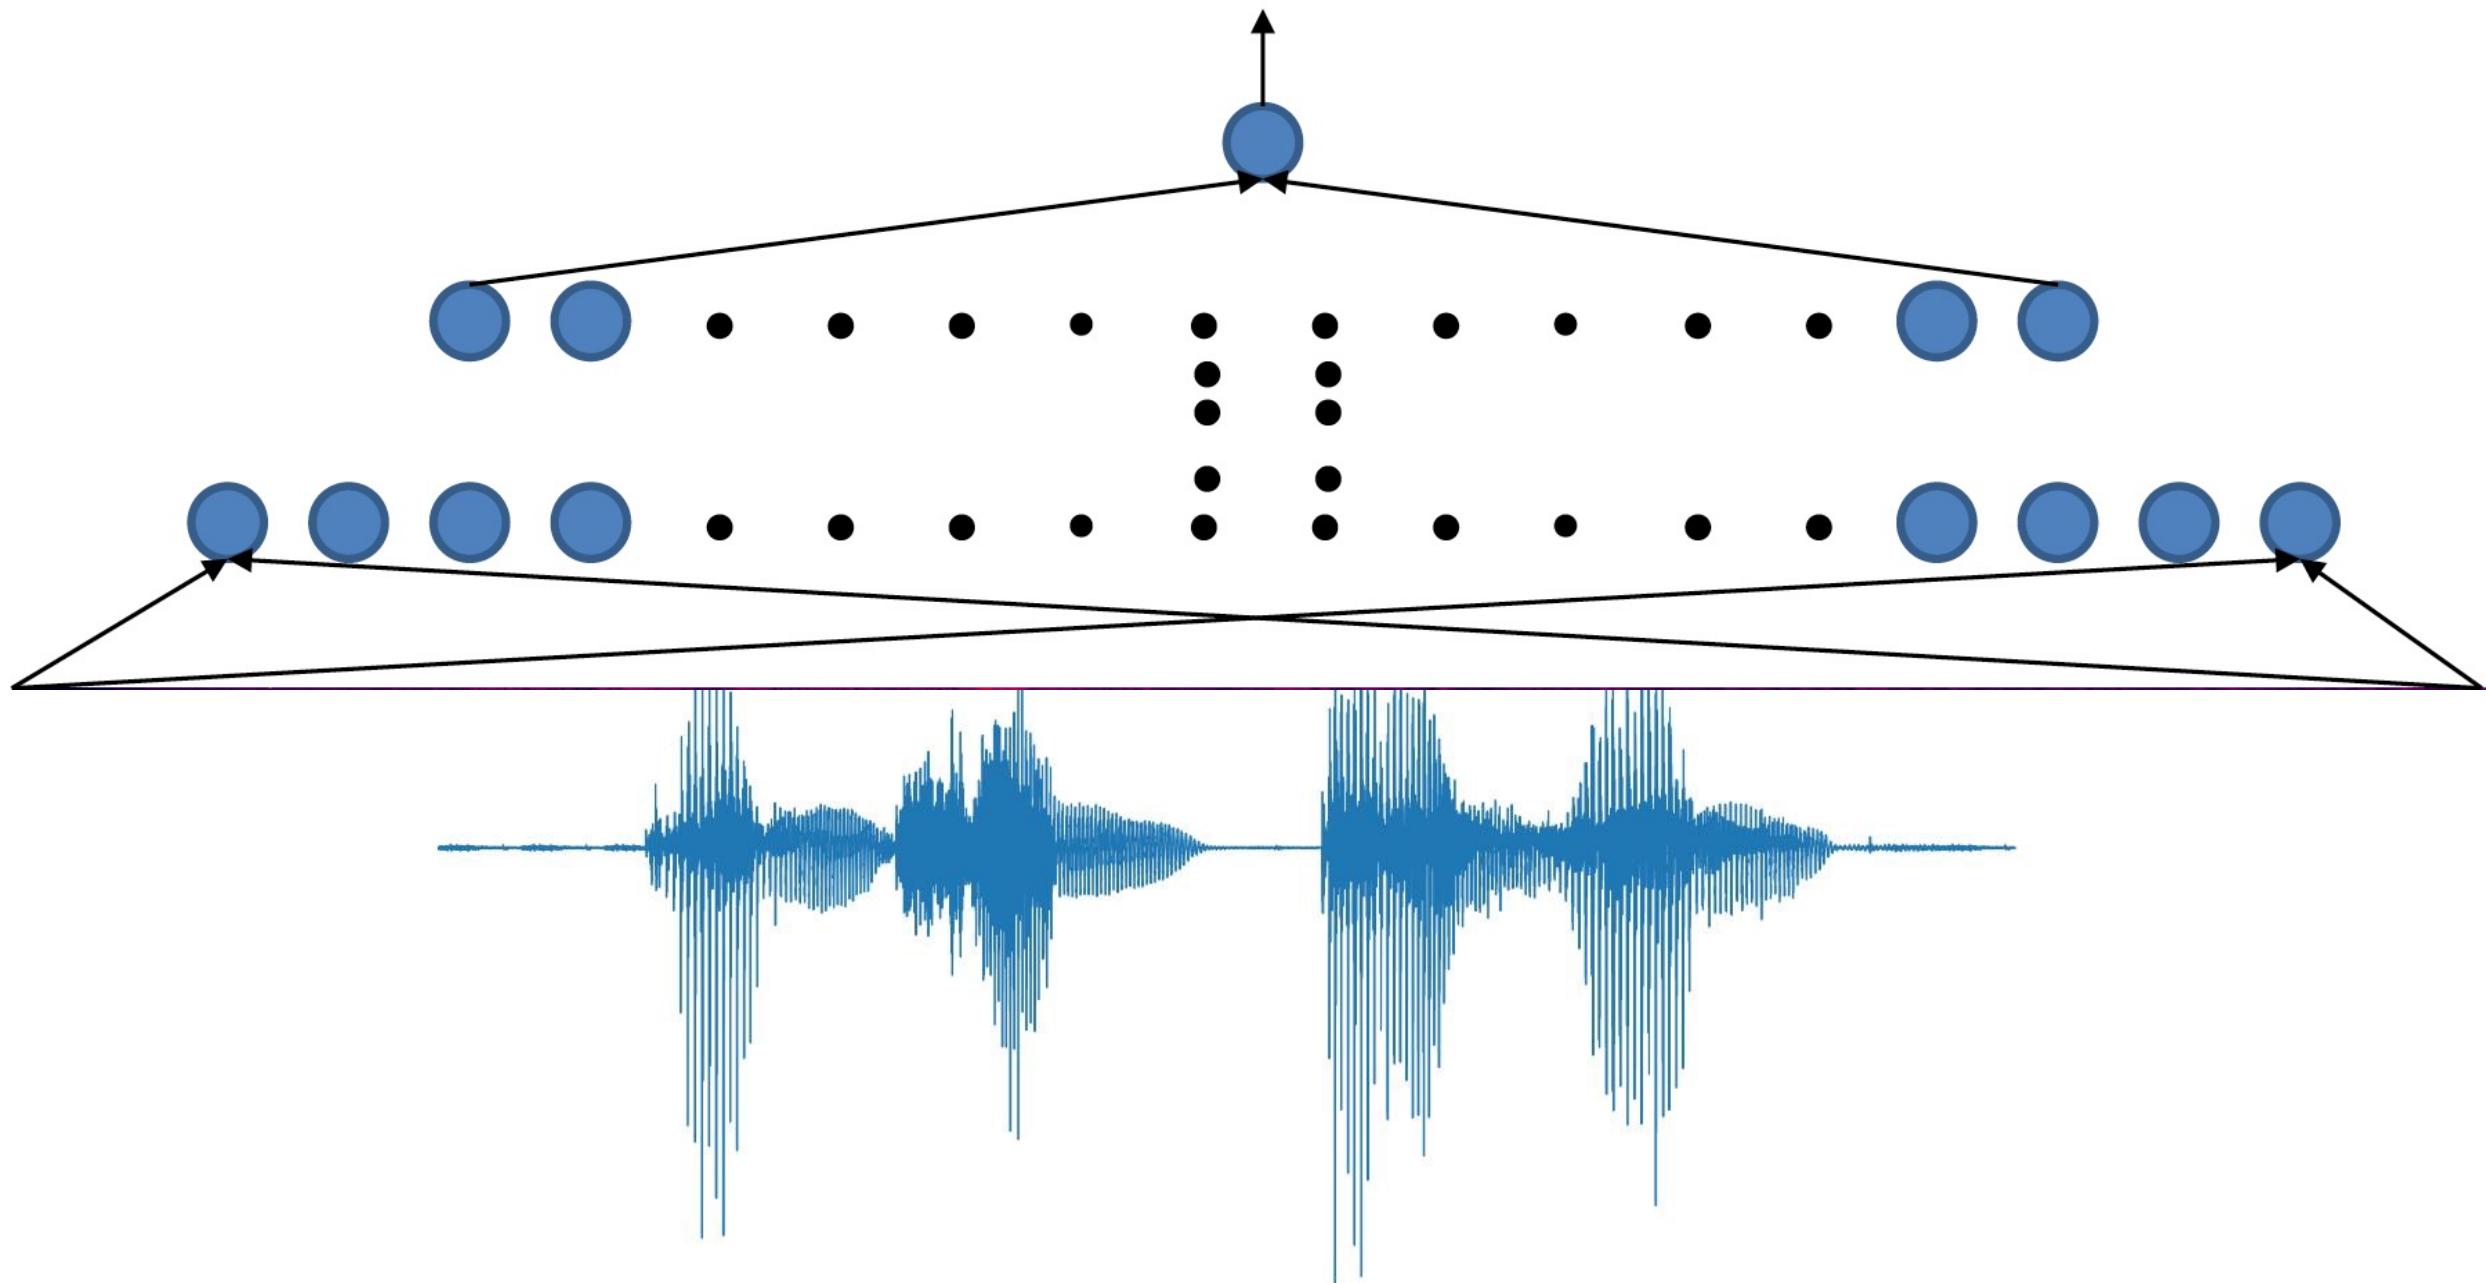
Search Pattern

Search Pattern



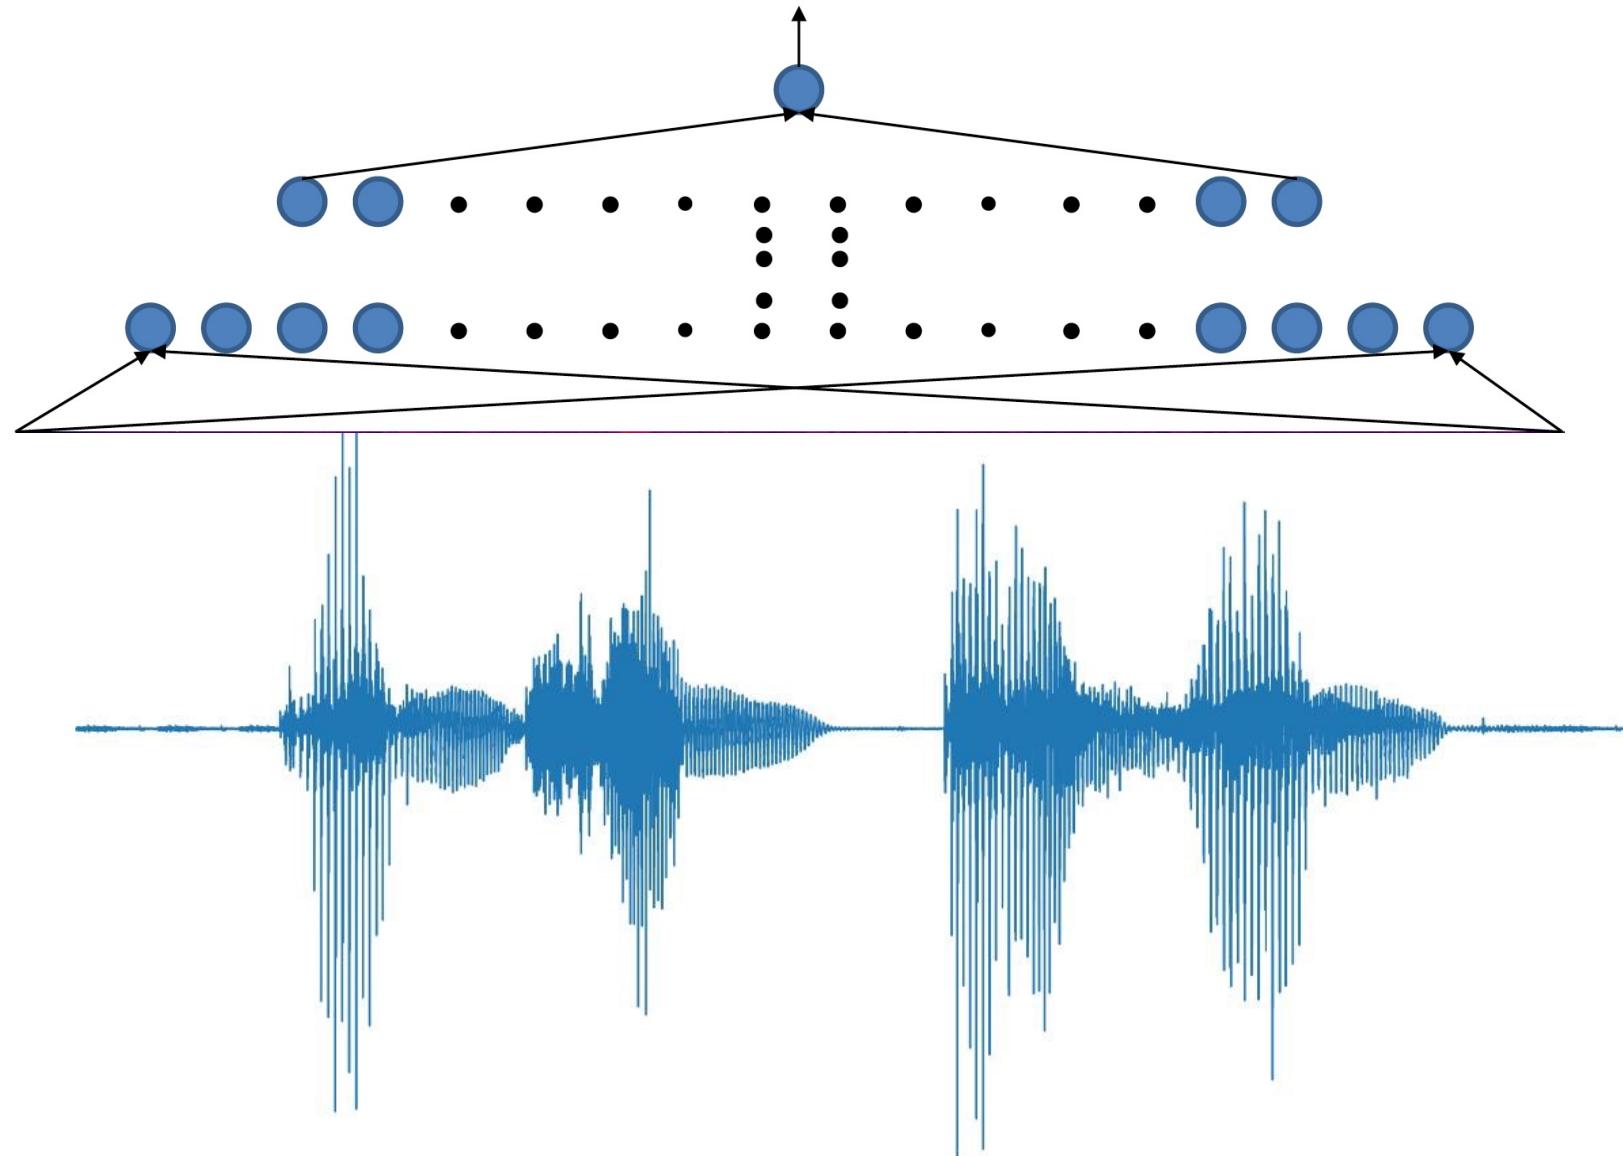
Using MLP to answer:

- Does this signal contain the word “Machine Learning”?
- If yes, where is the word “Machine Learning” ?



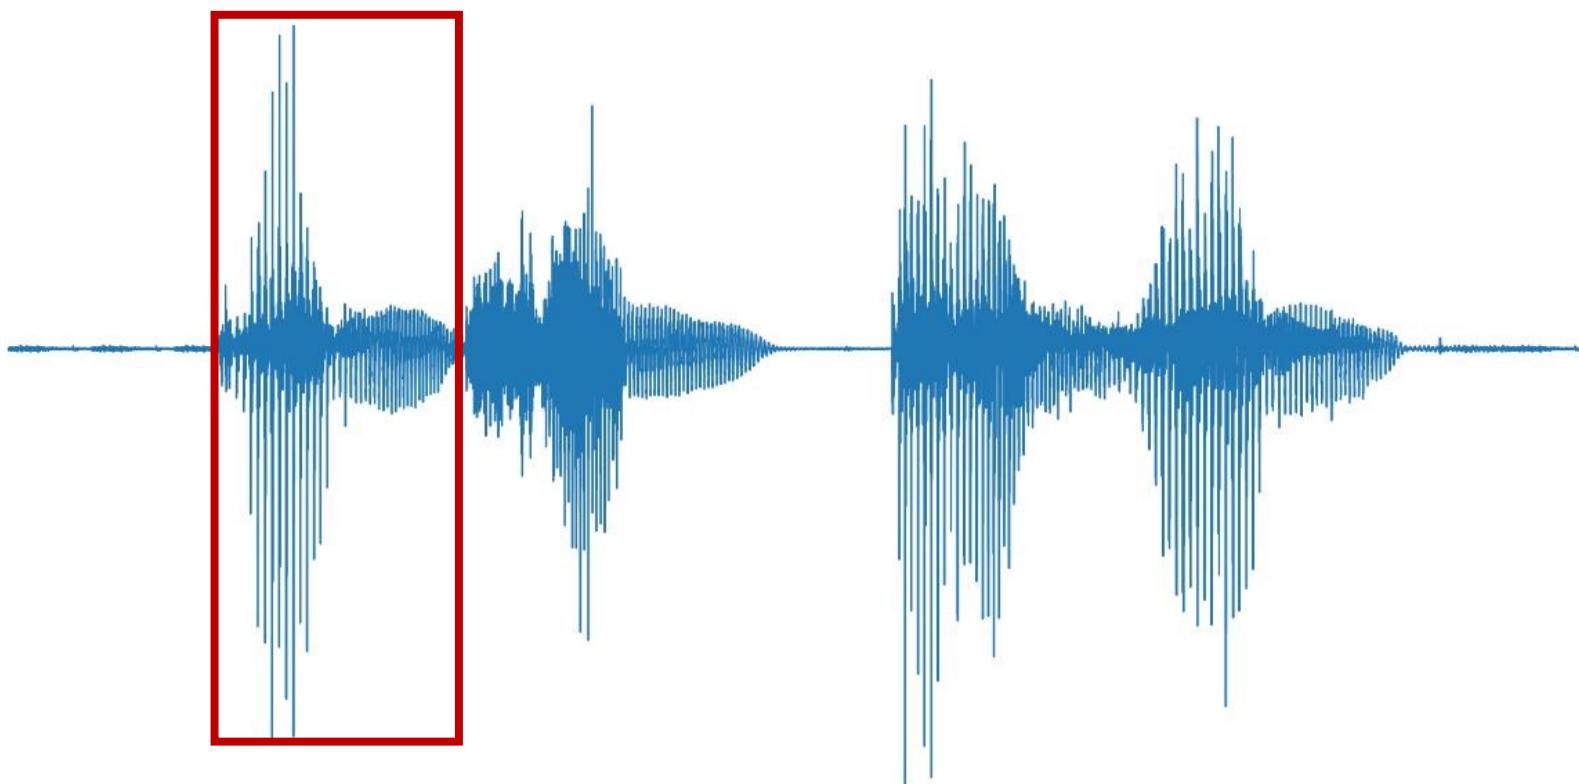
Search Pattern with MLP

Problems?

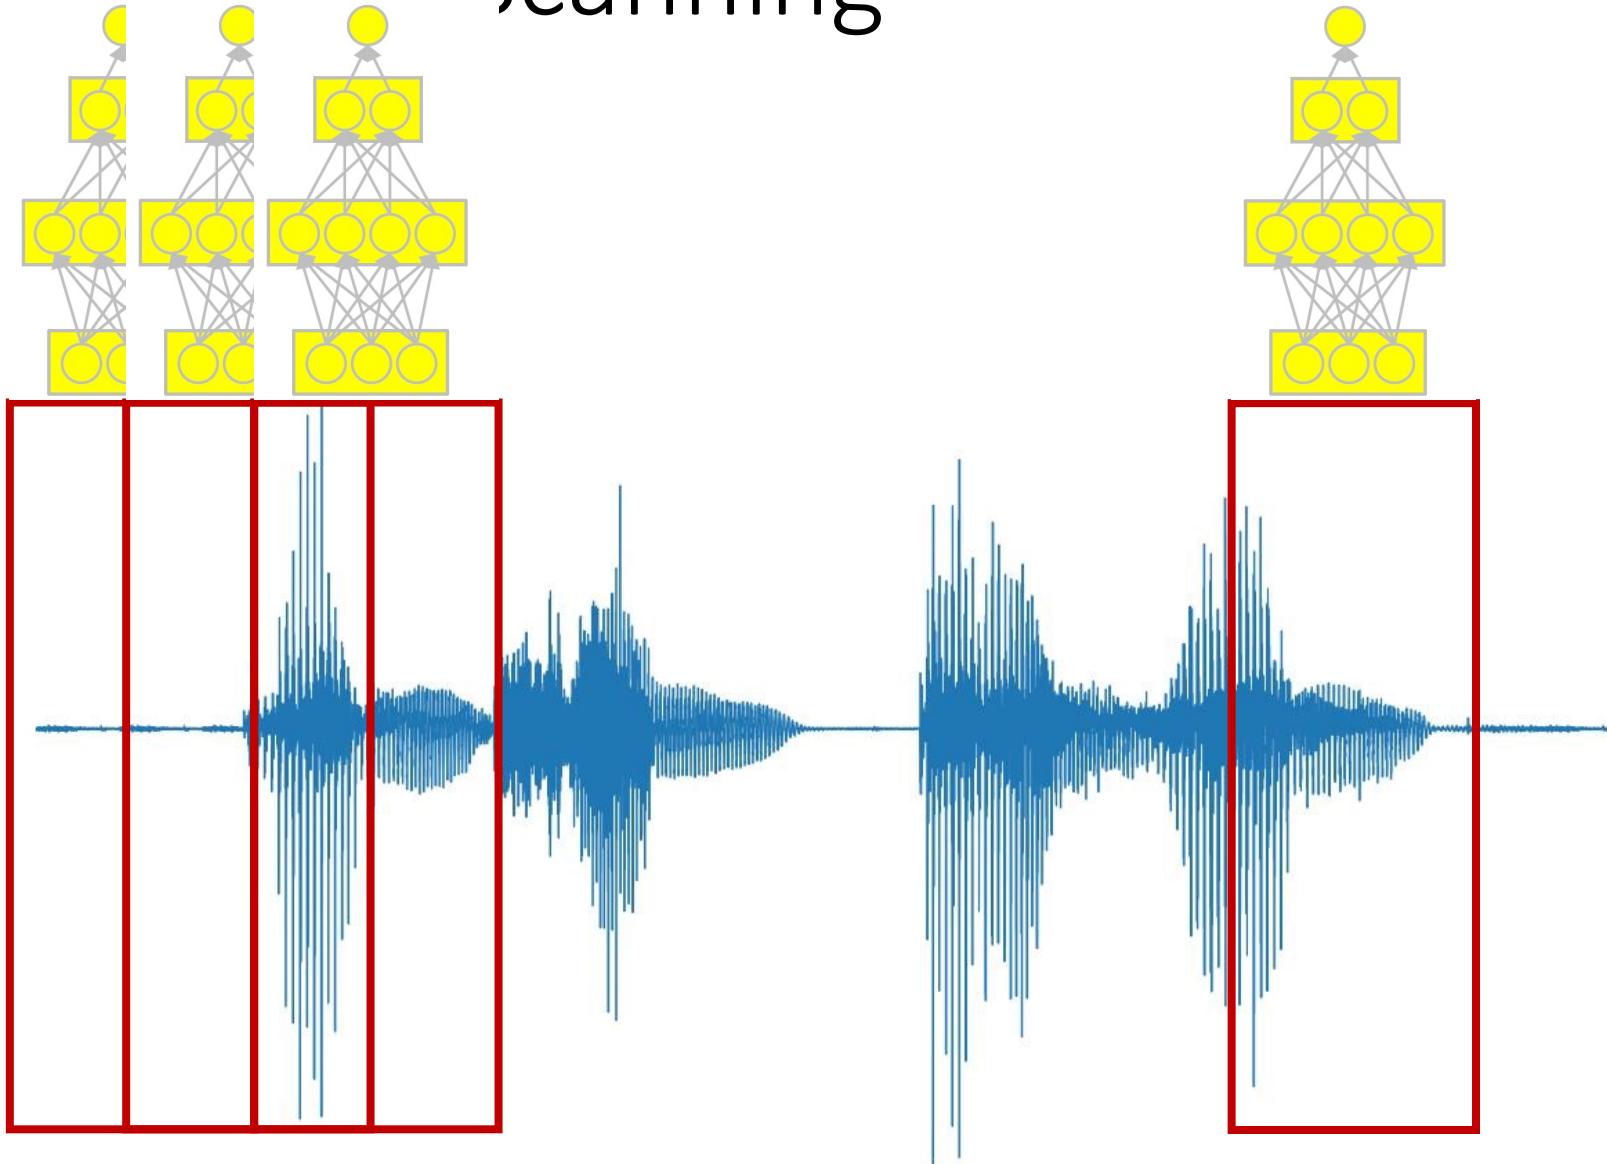


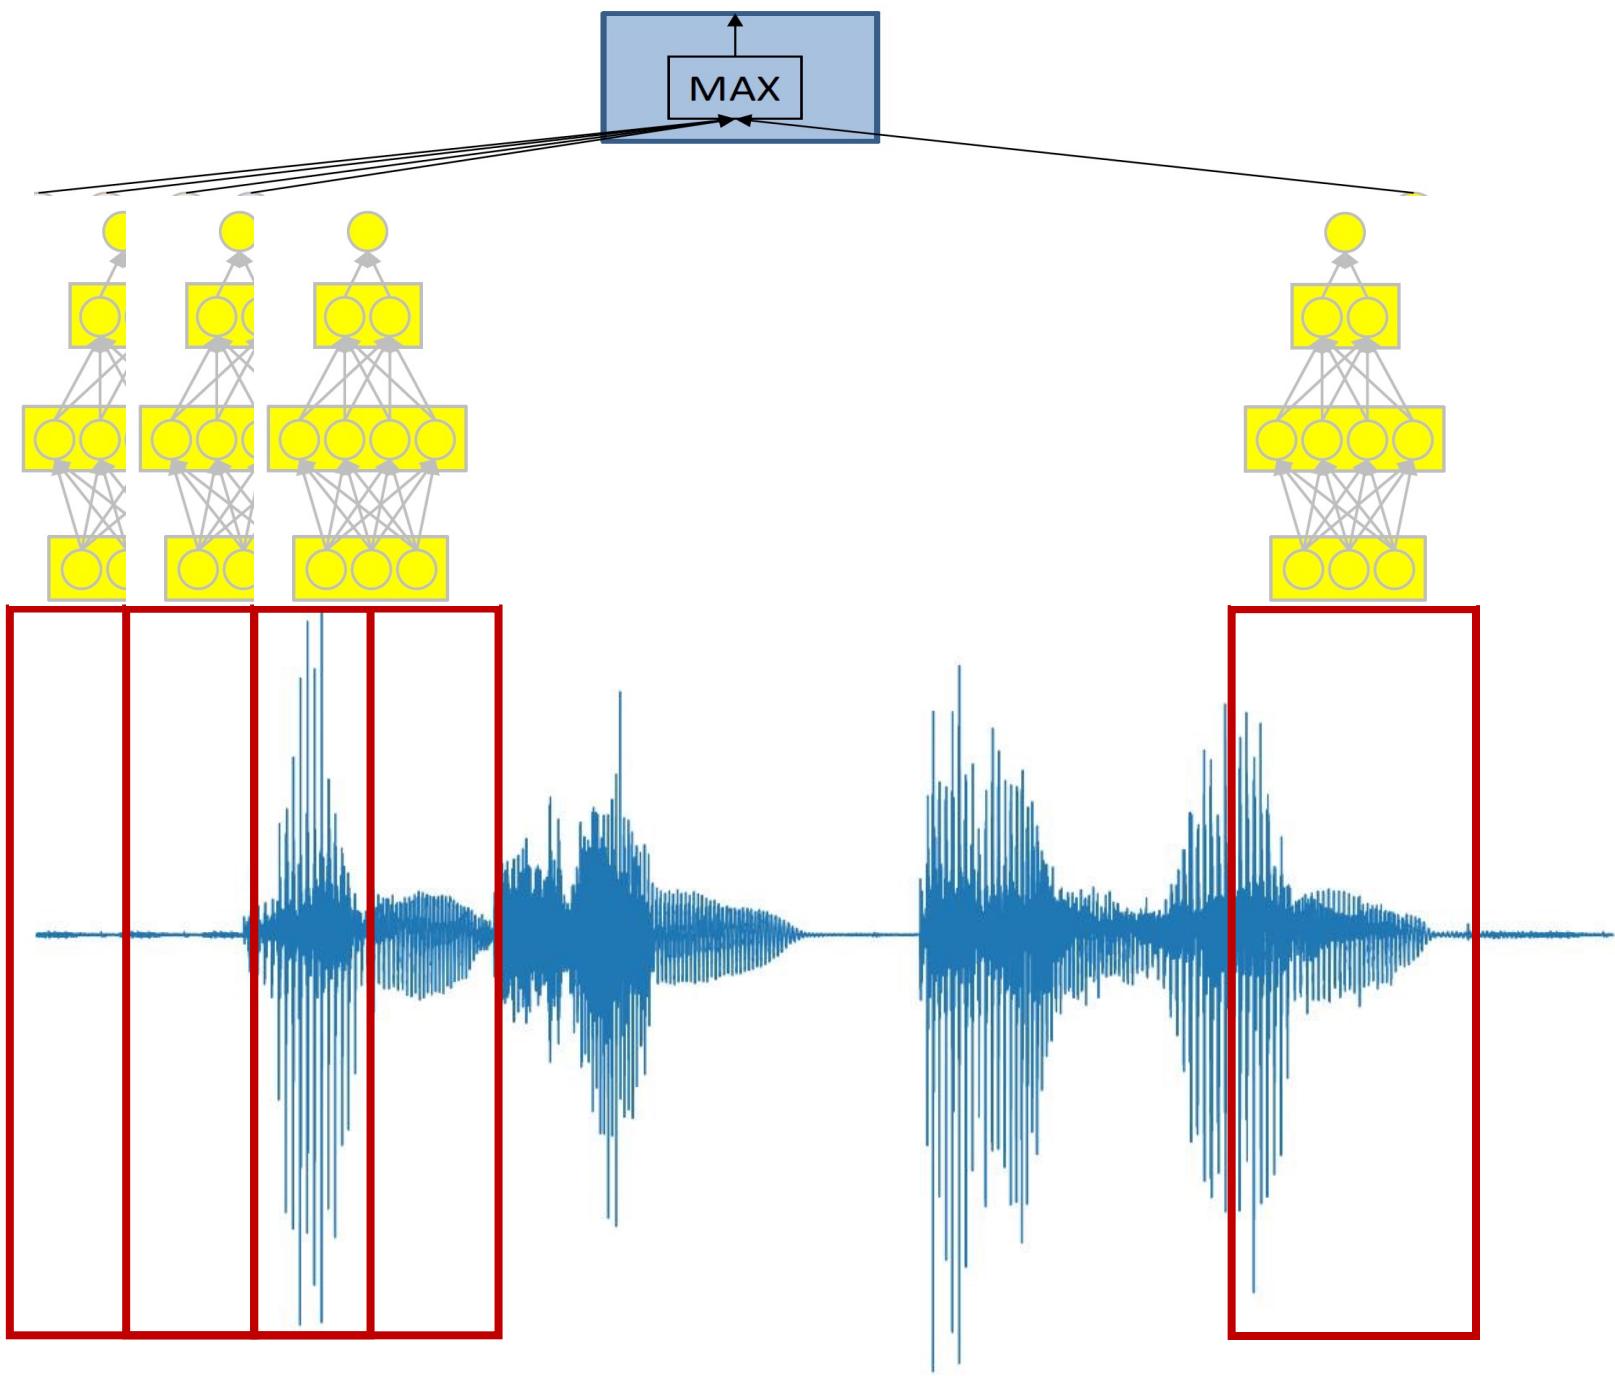
Search Pattern

Scanning



Search Pattern Scanning





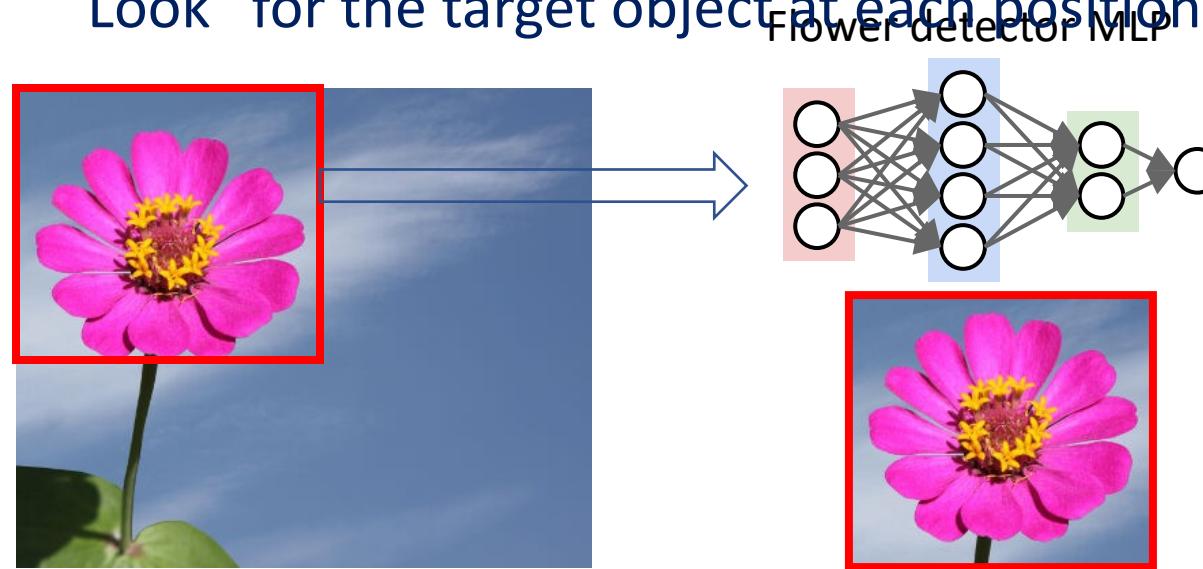
Flower Detection by scanning

“Look” for the target object at each position



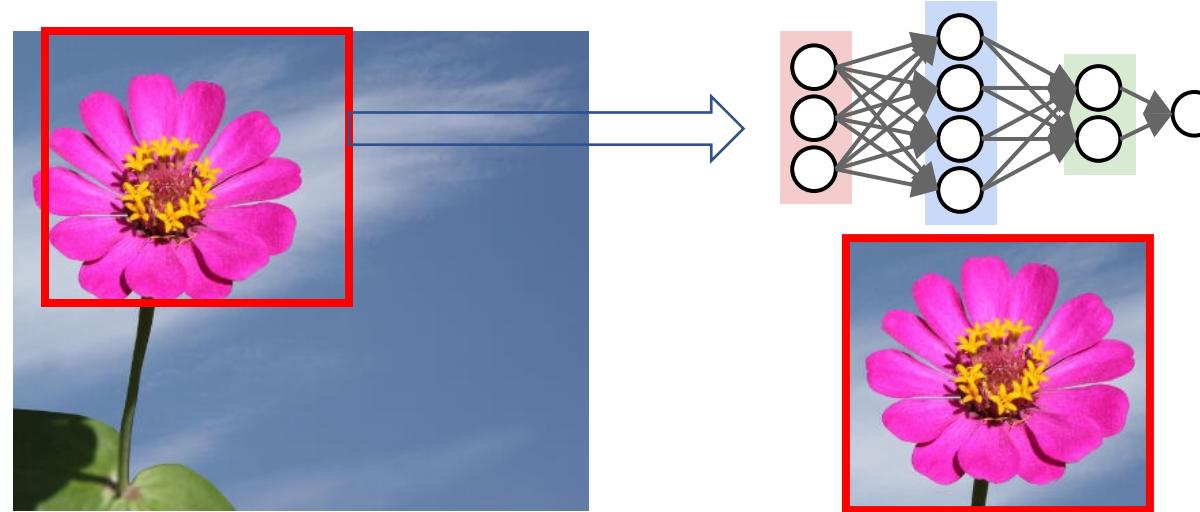
Flower Detection by scanning

“Look” for the target object at each position

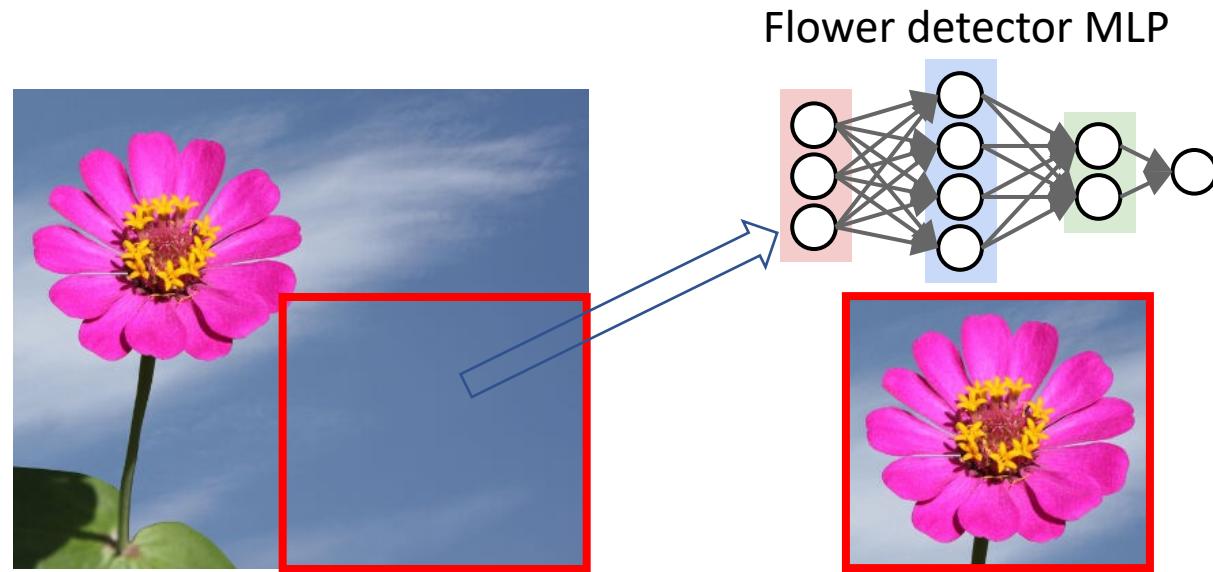


Flower Detection by scanning

“Look” for the target object at each position

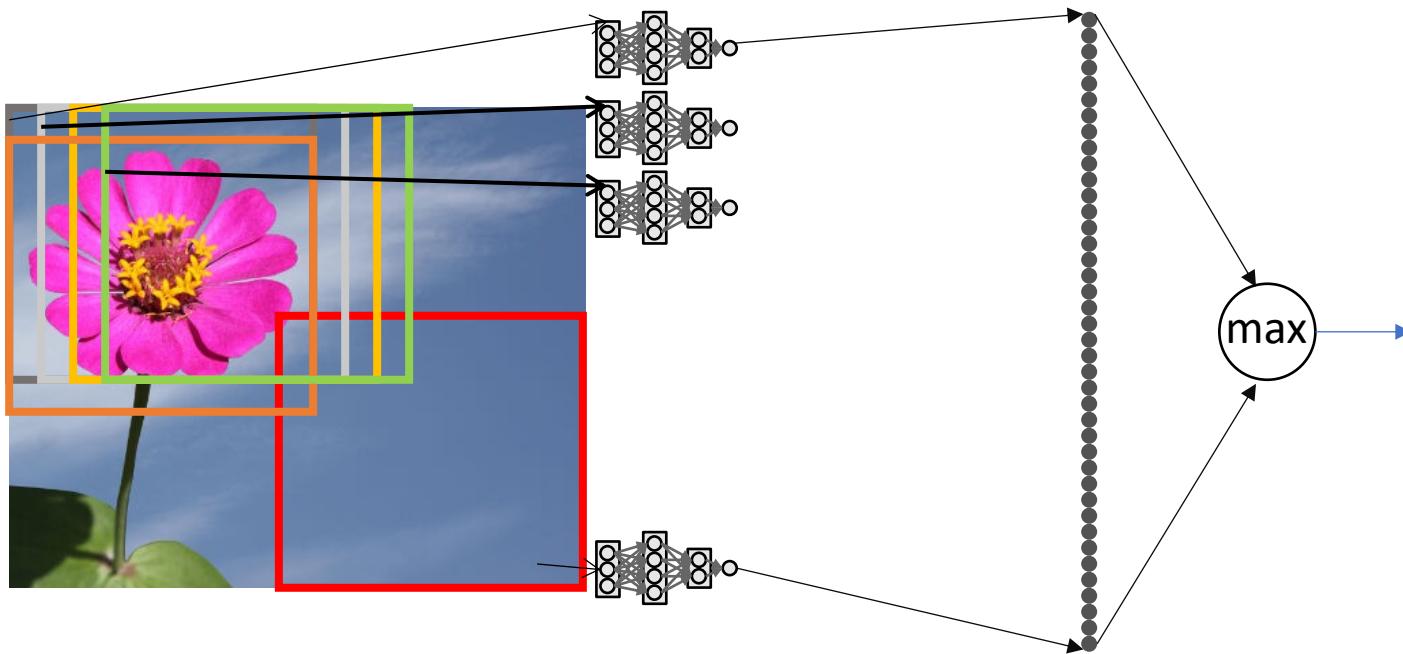


Flower Detection by scanning

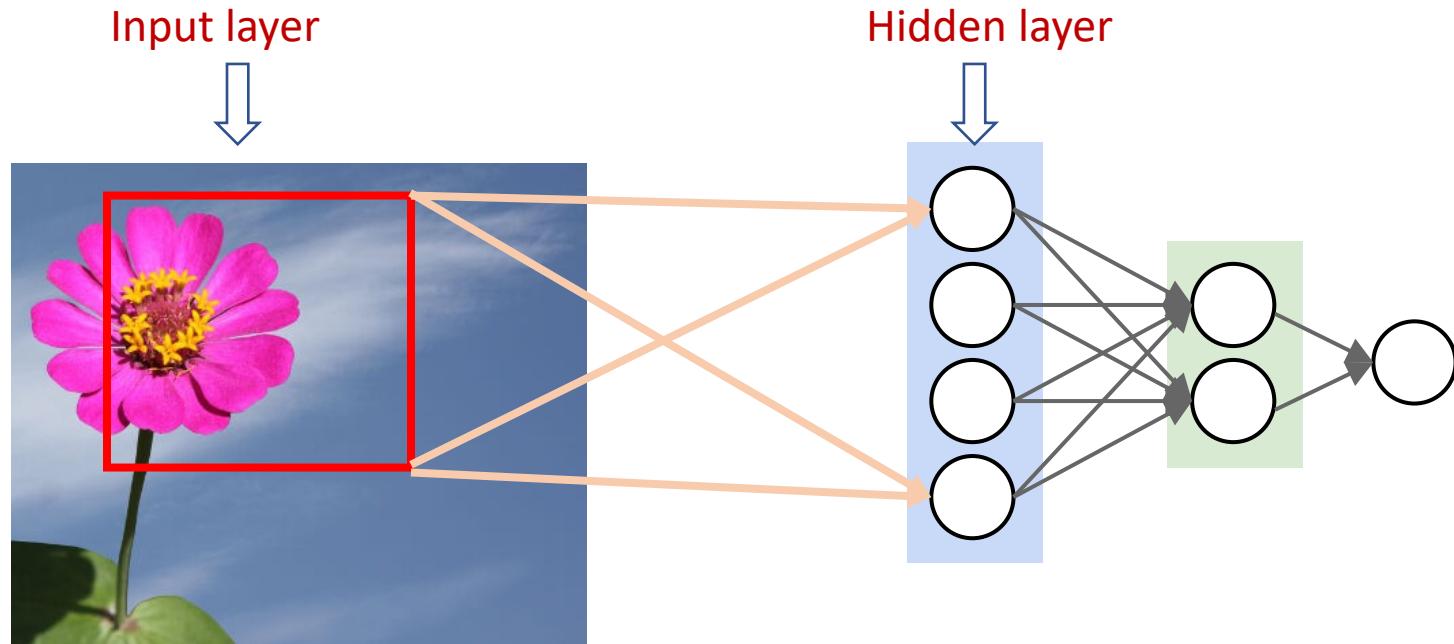


Flower Detection by scanning

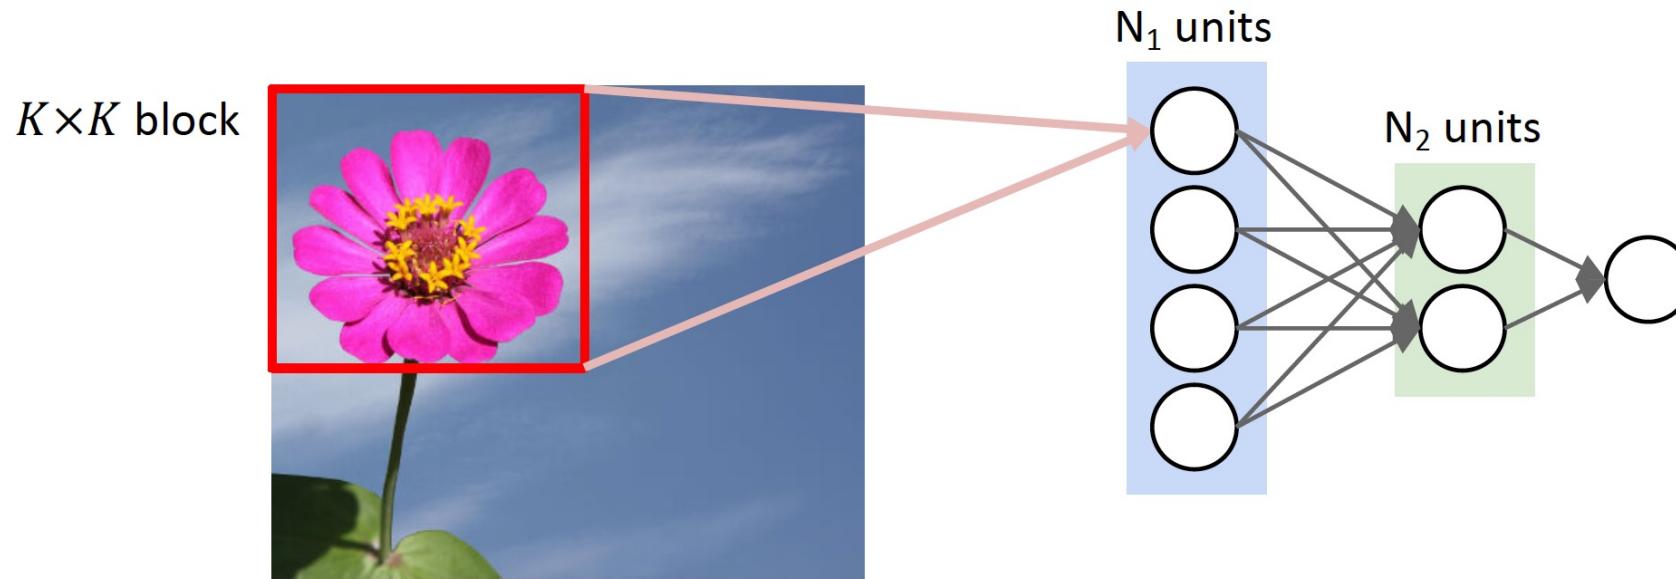
“Look” for the target object at each position



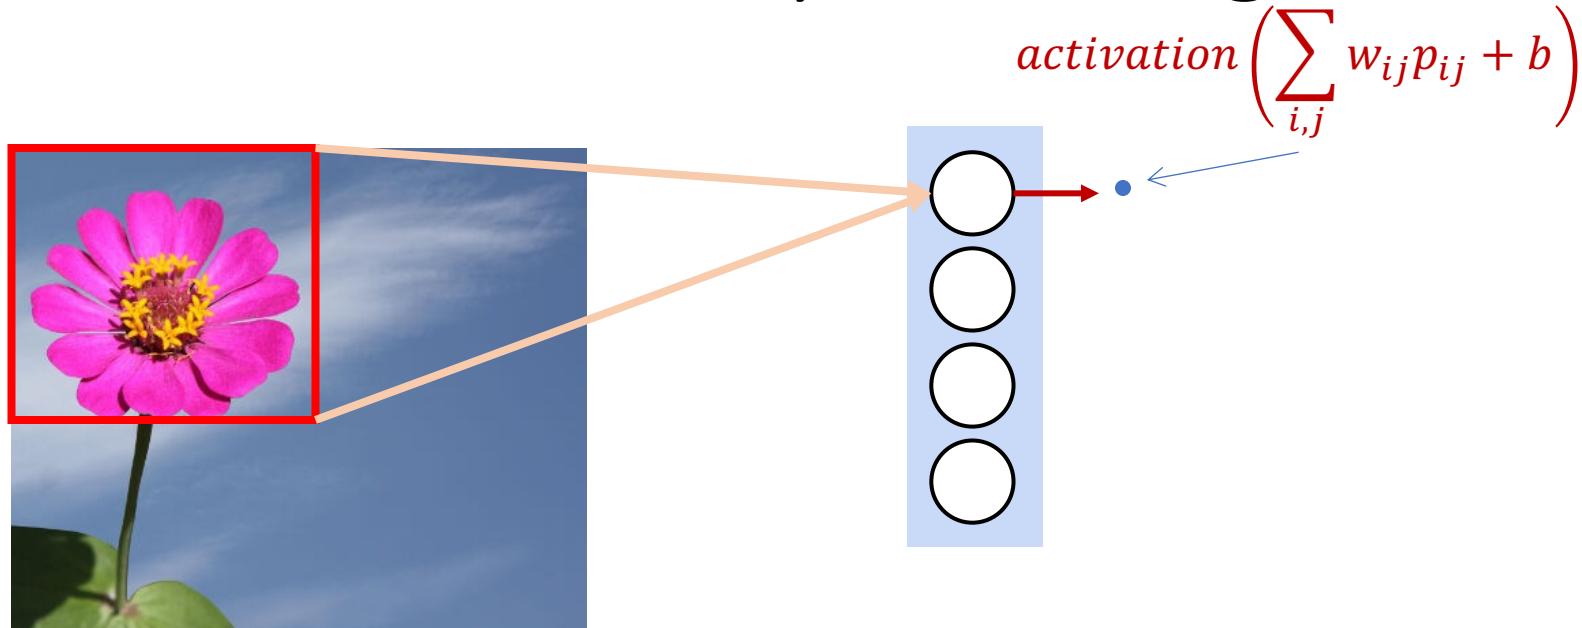
Flower Detection by scanning



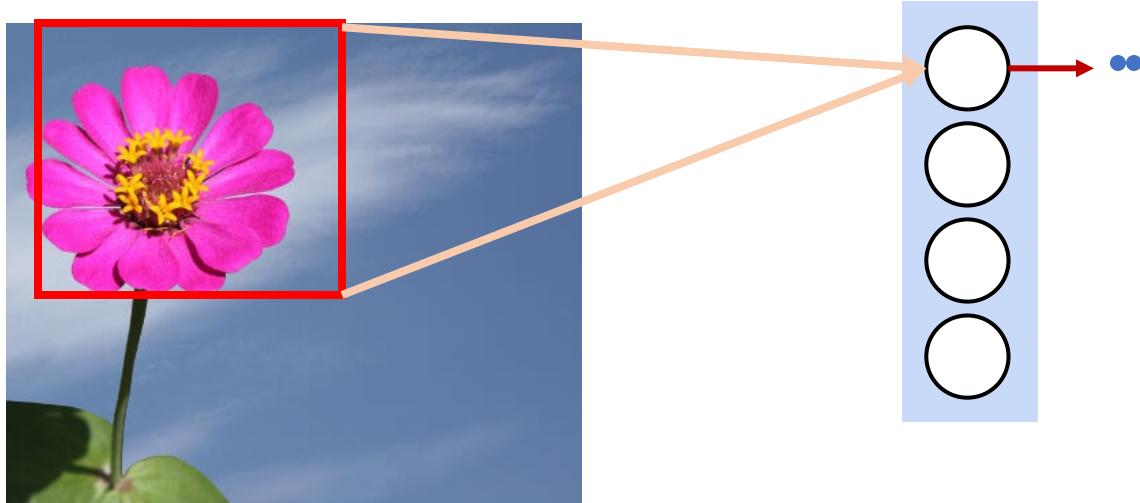
Flower Detection by scanning



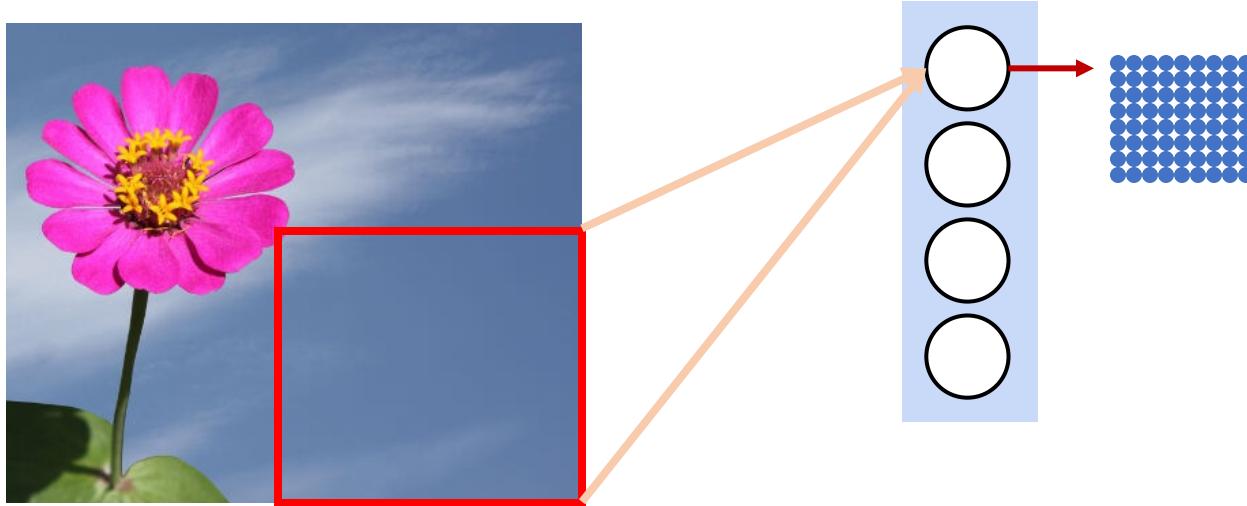
Flower Detection by scanning



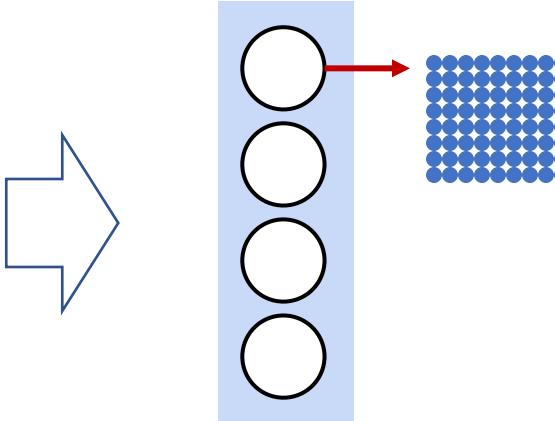
Flower Detection by scanning



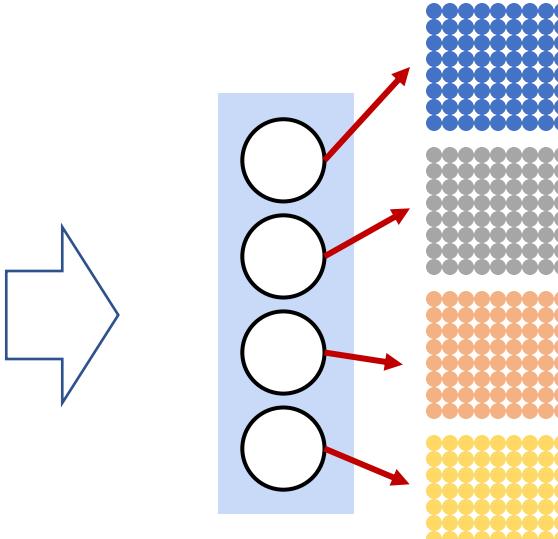
Flower Detection by scanning



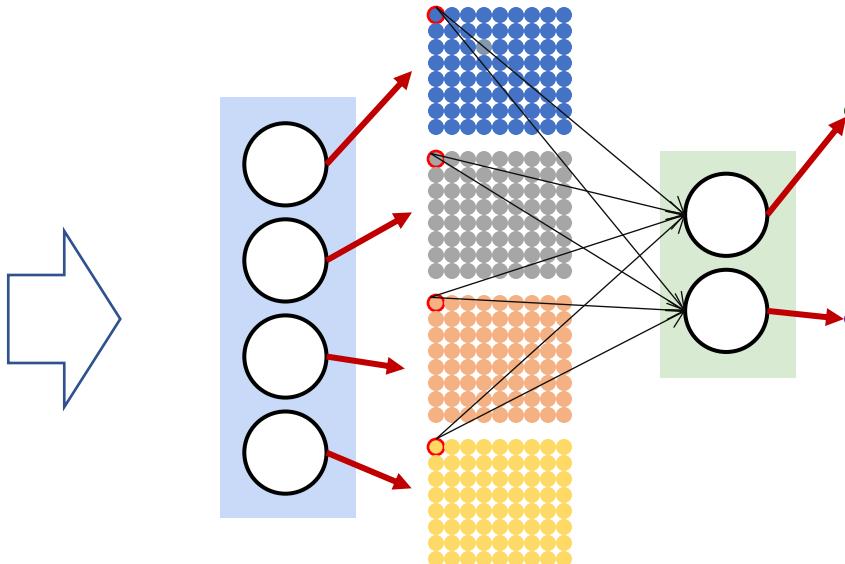
Flower Detection by scanning



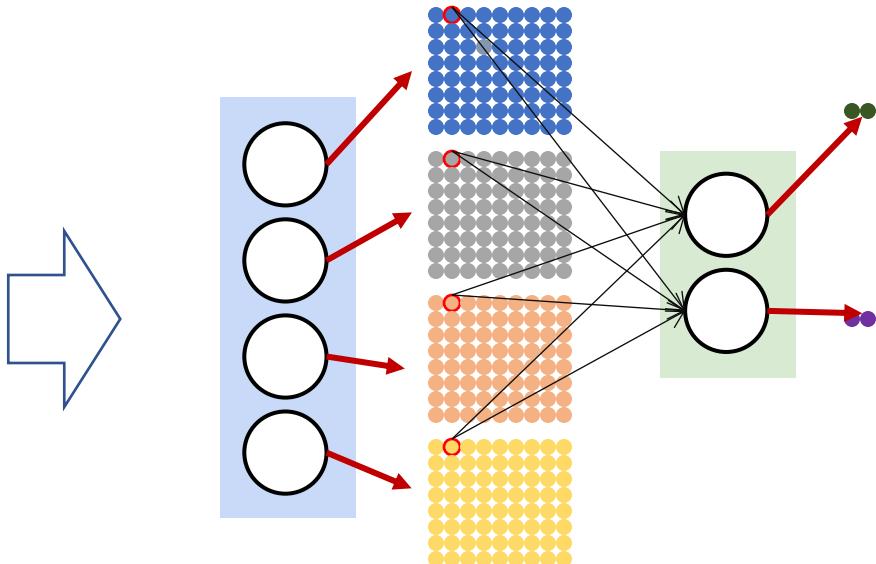
Flower Detection by scanning



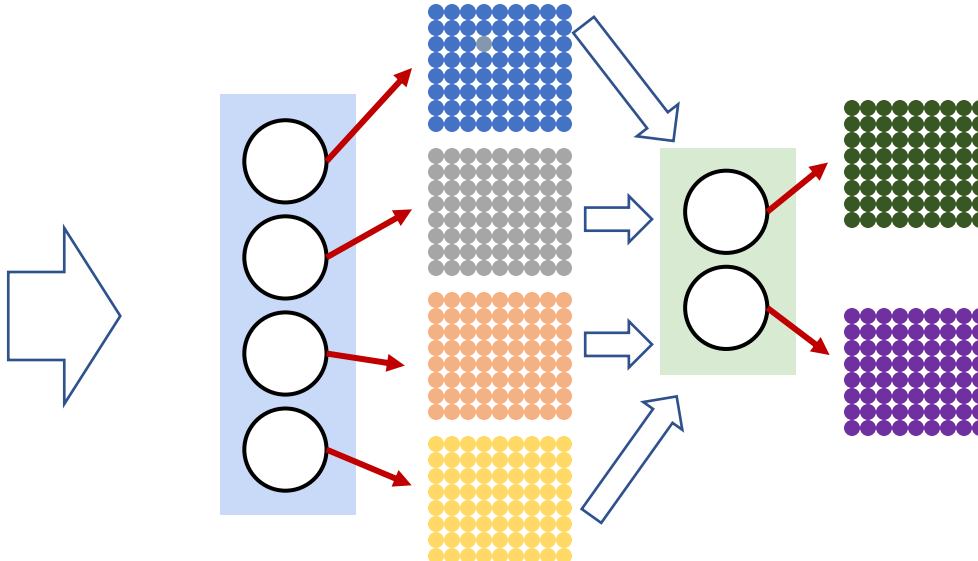
Flower Detection by scanning



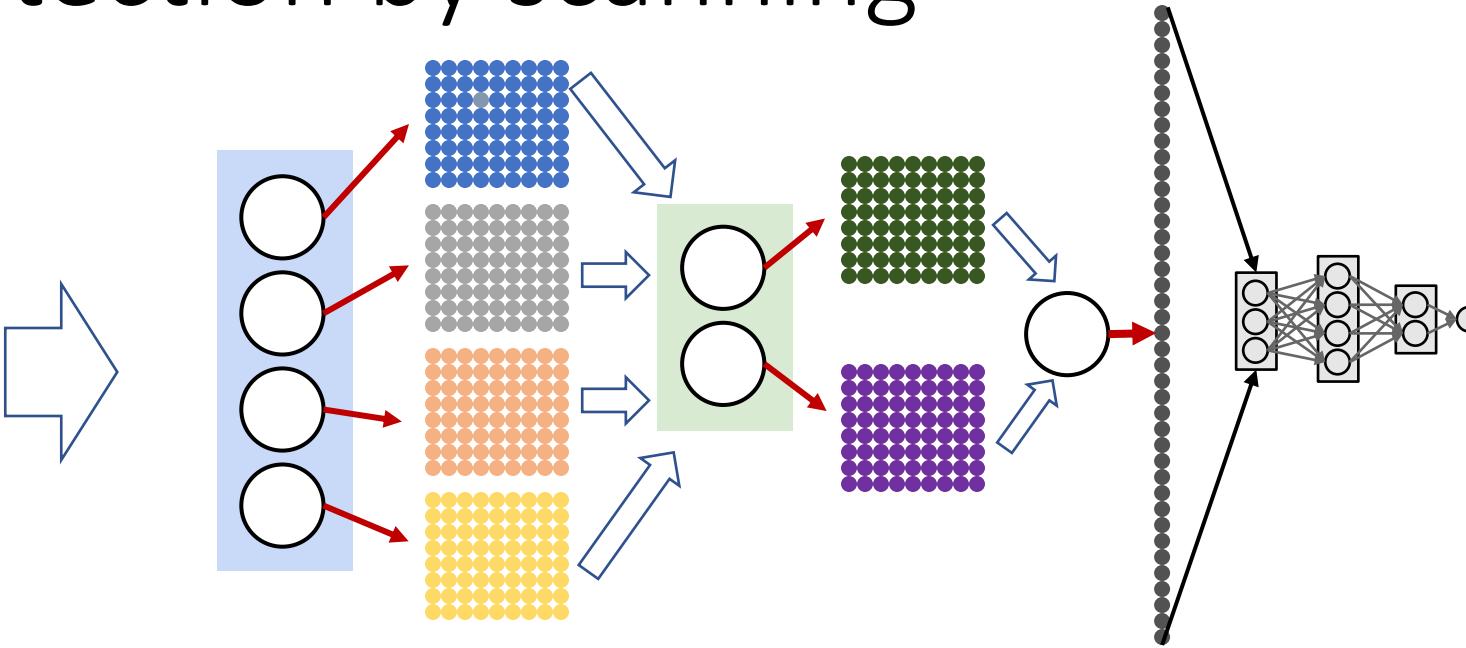
Flower Detection by scanning



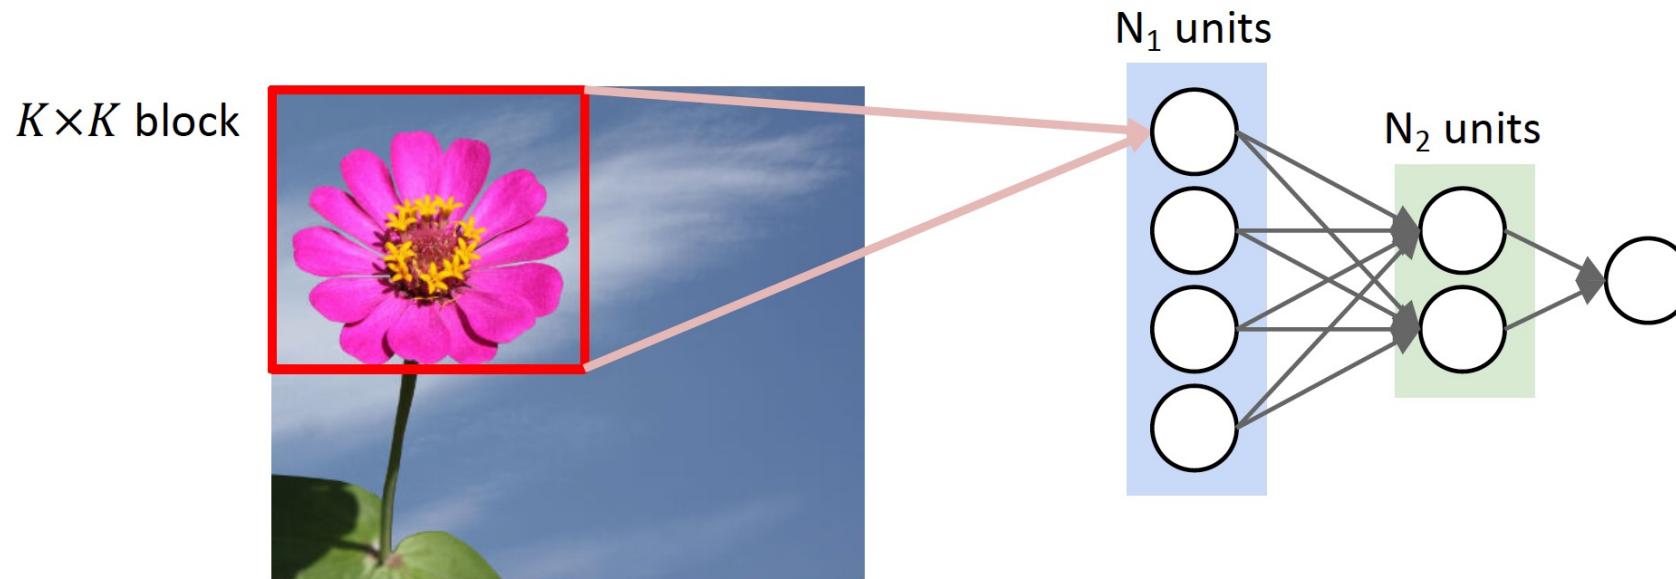
Flower Detection by scanning



Flower Detection by scanning



Flower Detection by scanning



$(K^2 + 1)N_1$ weights in first layer

$(N_1 + 1)N_2$ weights in second layer

$(N_{i-1} + 1)N_i$ weights in subsequent i^{th} layer

Total parameters:

$$\mathcal{O}(K^2 N_1 + N_1 N_2 + N_2 N_3 \dots)$$

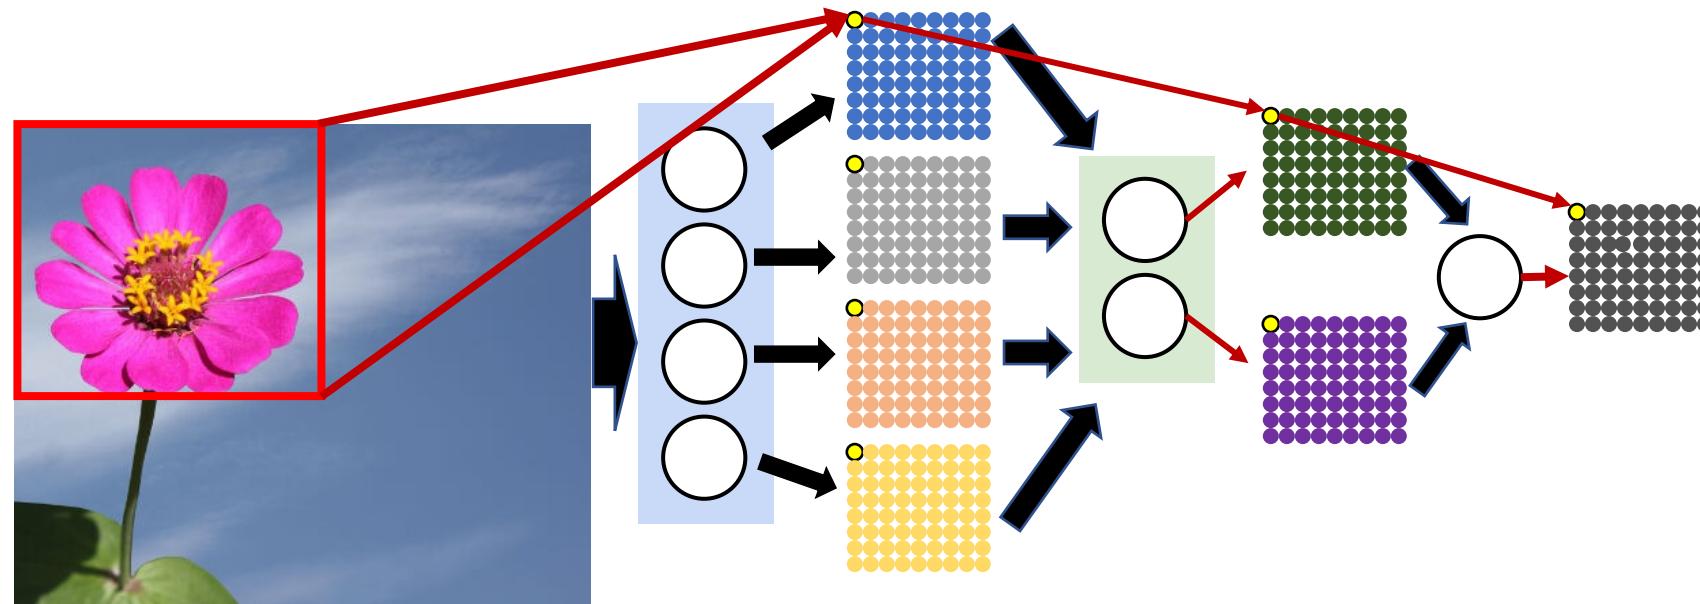
- Example

For this example, let $K = 16, N_1 = 4, N_2 = 2, N_3 = 1$

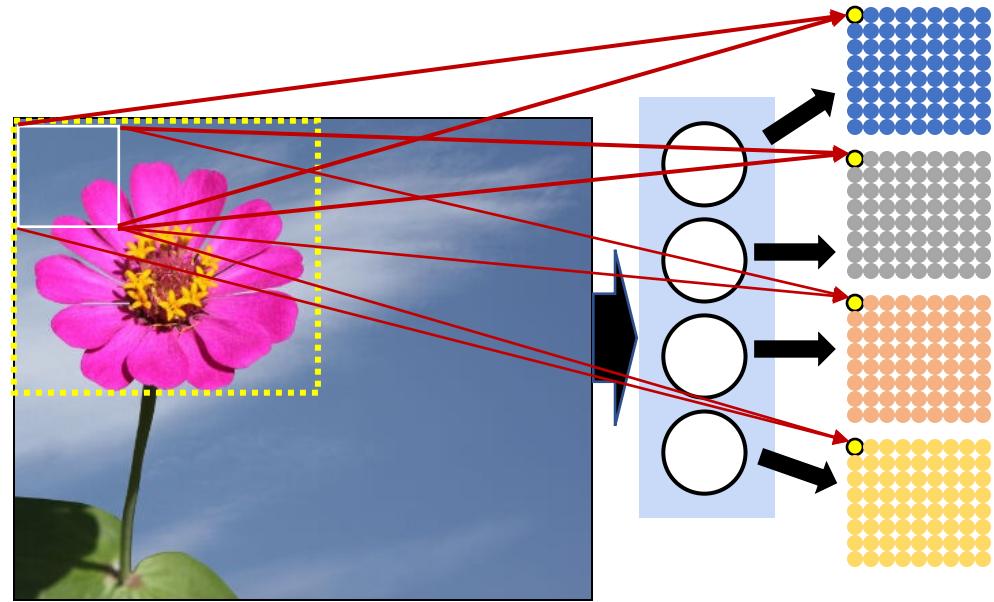
Total 1034 weights

Scanning with Distribution

Flower Detection by scanning



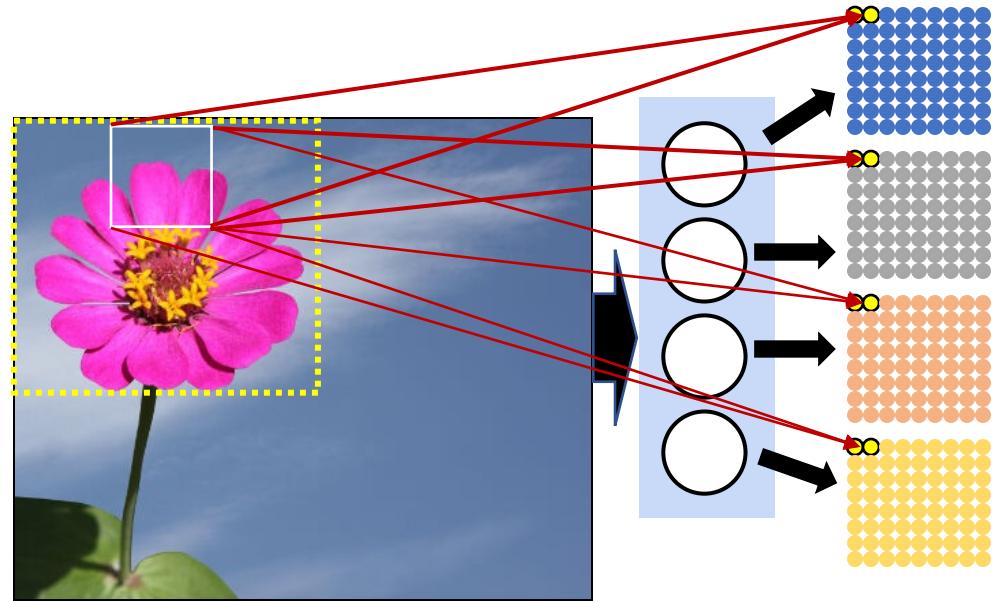
Flower Detection: Distributing the scan



The first layer evaluates smaller blocks of pixels

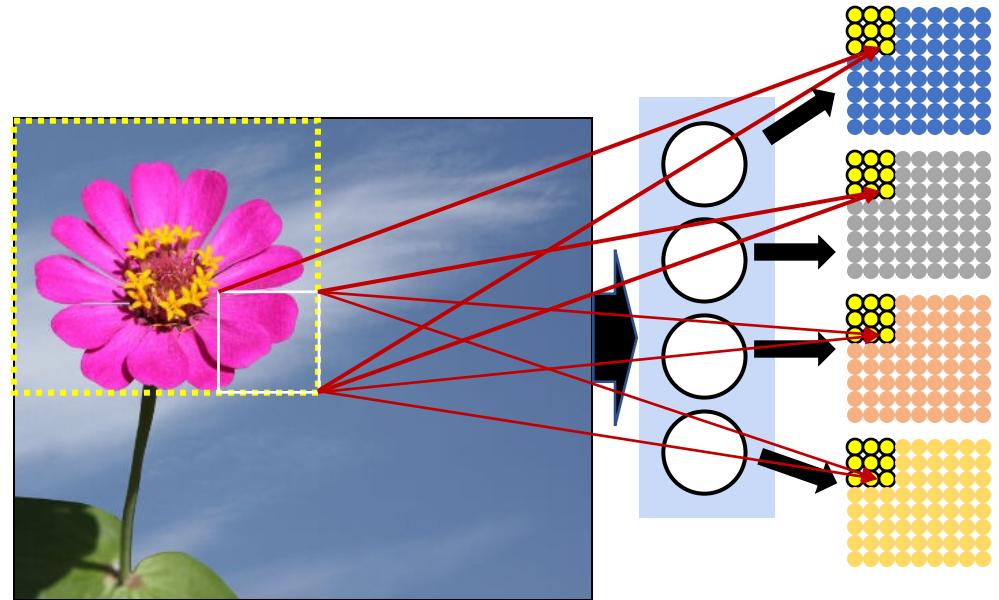
We can distribute the pattern matching over more (two) layers and still achieve the same block analysis at the second layer

Flower Detection: Distributing the scan



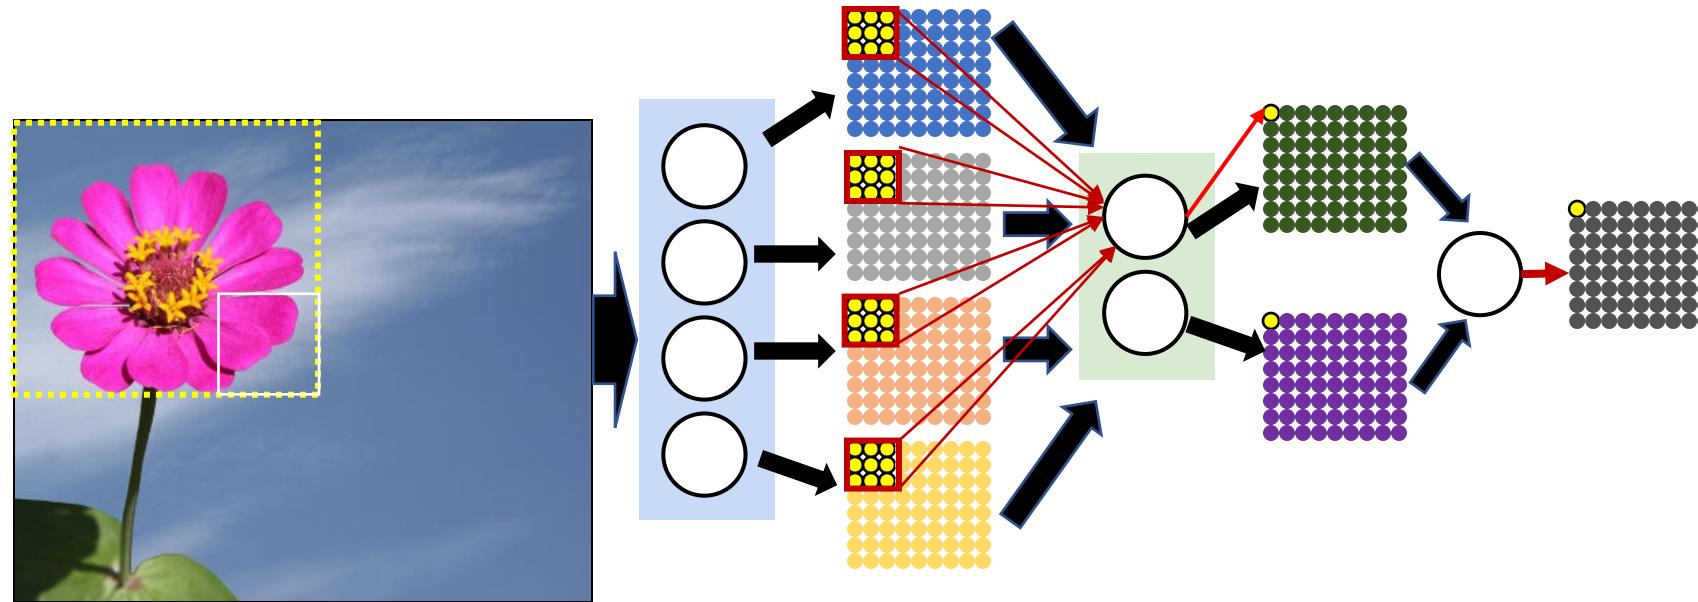
The first layer evaluates smaller blocks of pixels

Flower Detection: Distributing the scan



The first layer evaluates smaller blocks of pixels

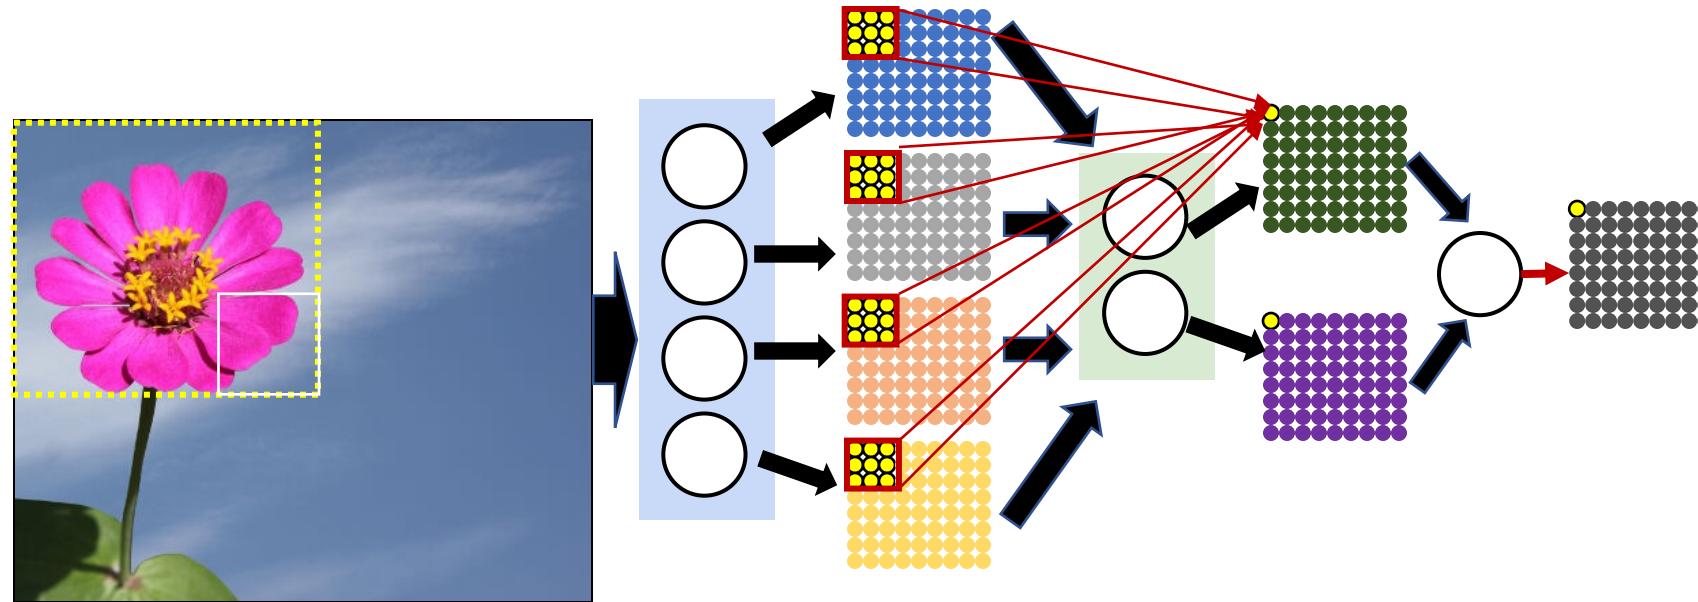
Flower Detection: Distributing the scan



The first layer evaluates smaller blocks of pixels

The next layer evaluates blocks of outputs from the first layer

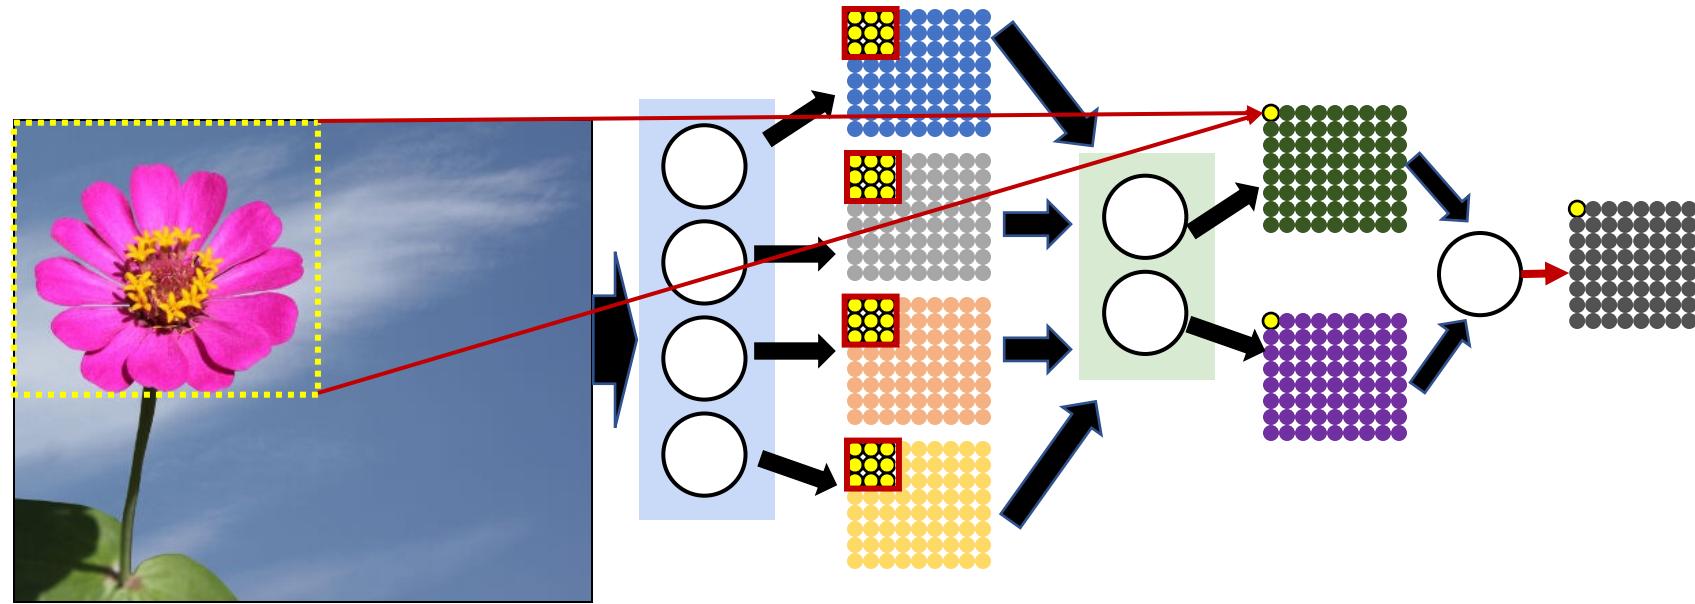
Flower Detection: Distributing the scan



The first layer evaluates smaller blocks of pixels

The next layer evaluates blocks of outputs from the first layer

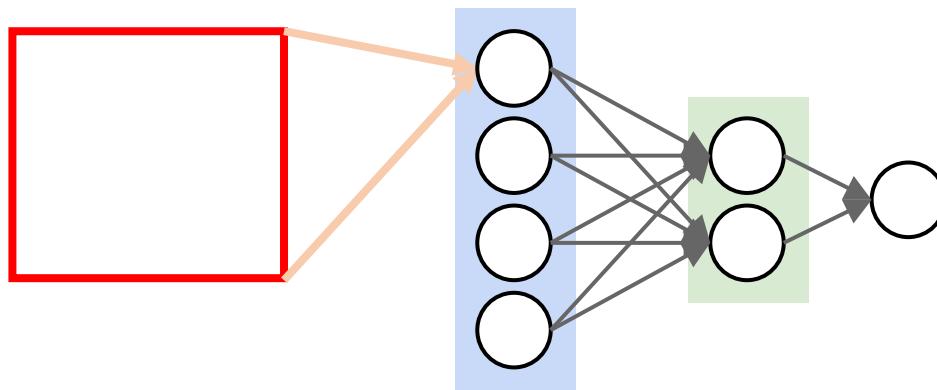
Flower Detection: Distributing the scan



The first layer evaluates smaller blocks of pixels

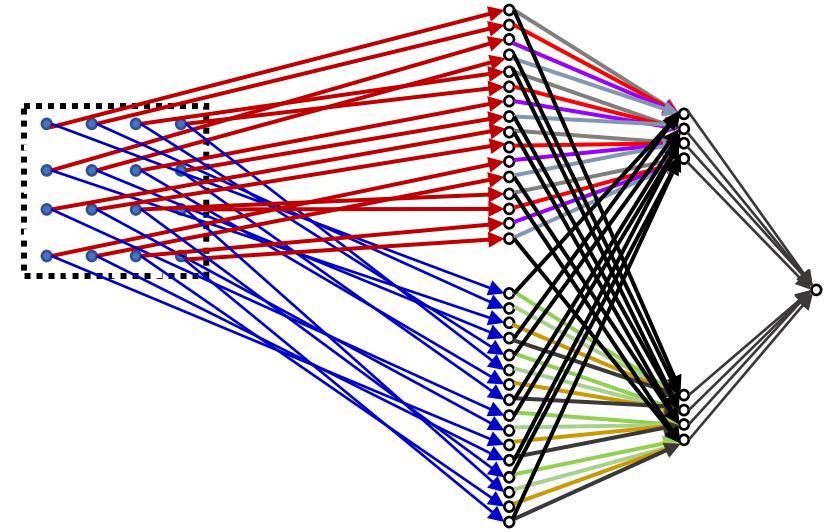
The next layer evaluates blocks of outputs from the first layer

Conventional MLP, not distributed

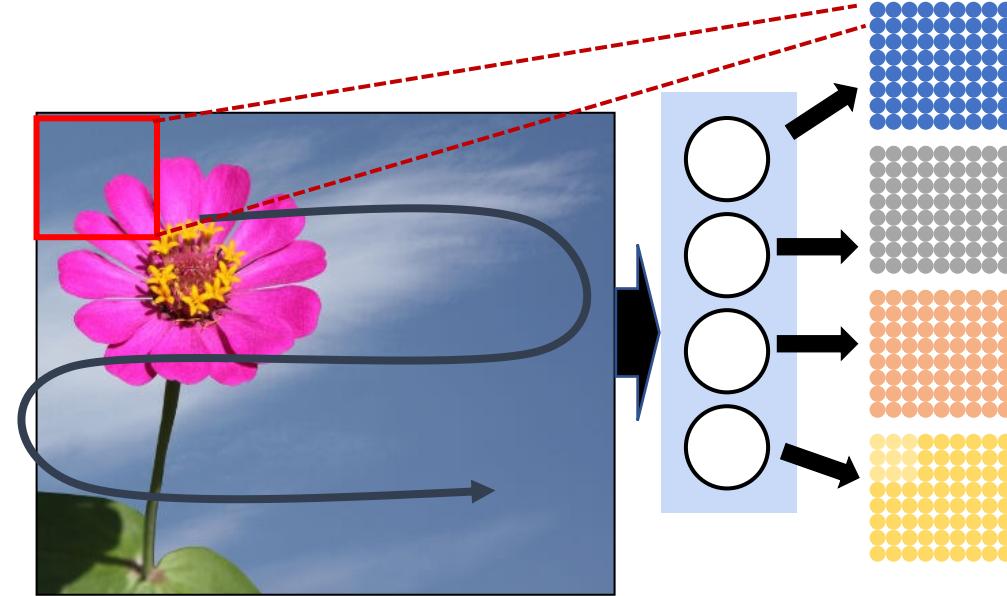


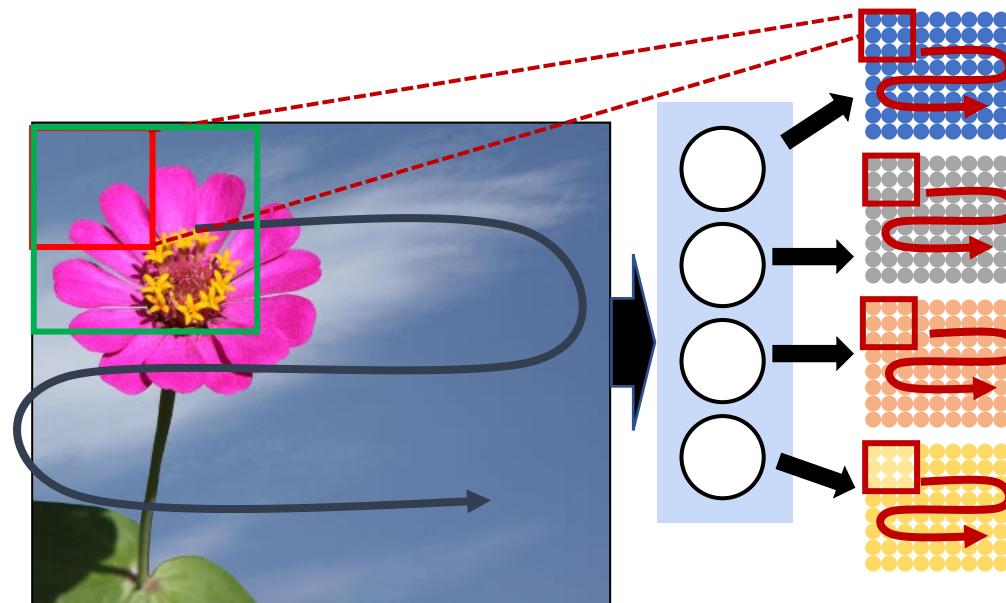
- $\mathcal{O}(K^2N_1 + N_1N_2 + N_2N_3 \dots)$
- For this example, let $K = 16, N_1 = 4, N_2 = 2, N_3 = 1$
- Total 1034 weights

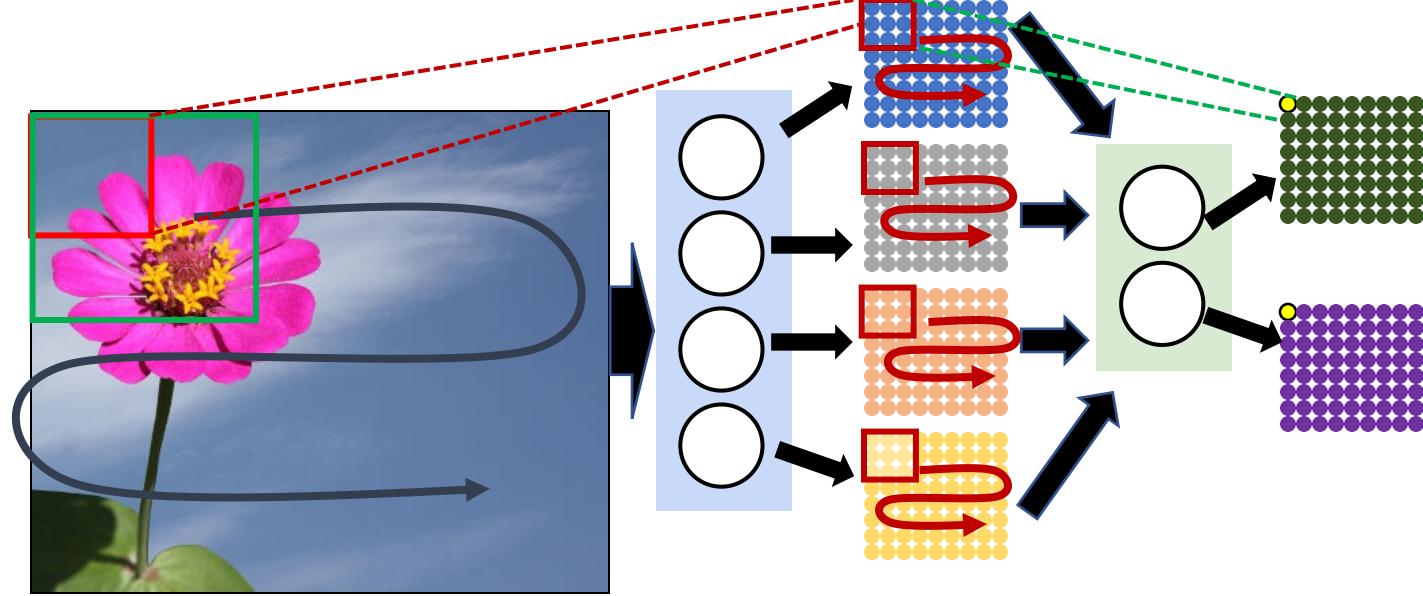
Distributed (3 layers)

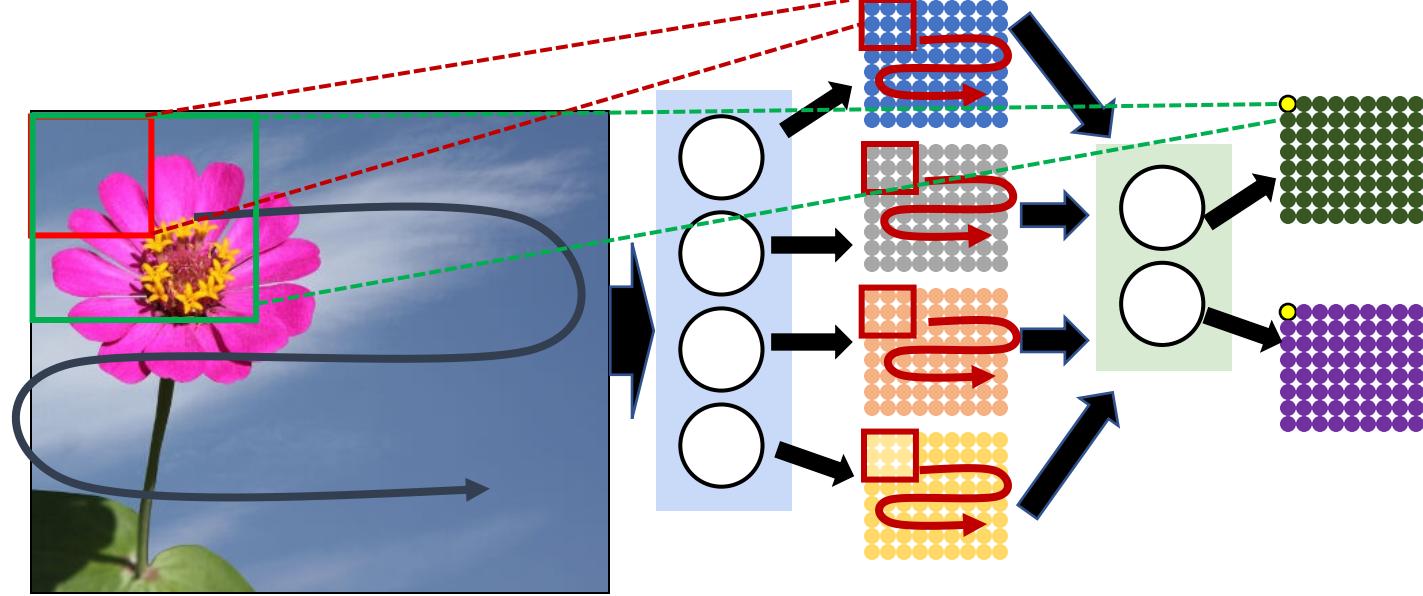


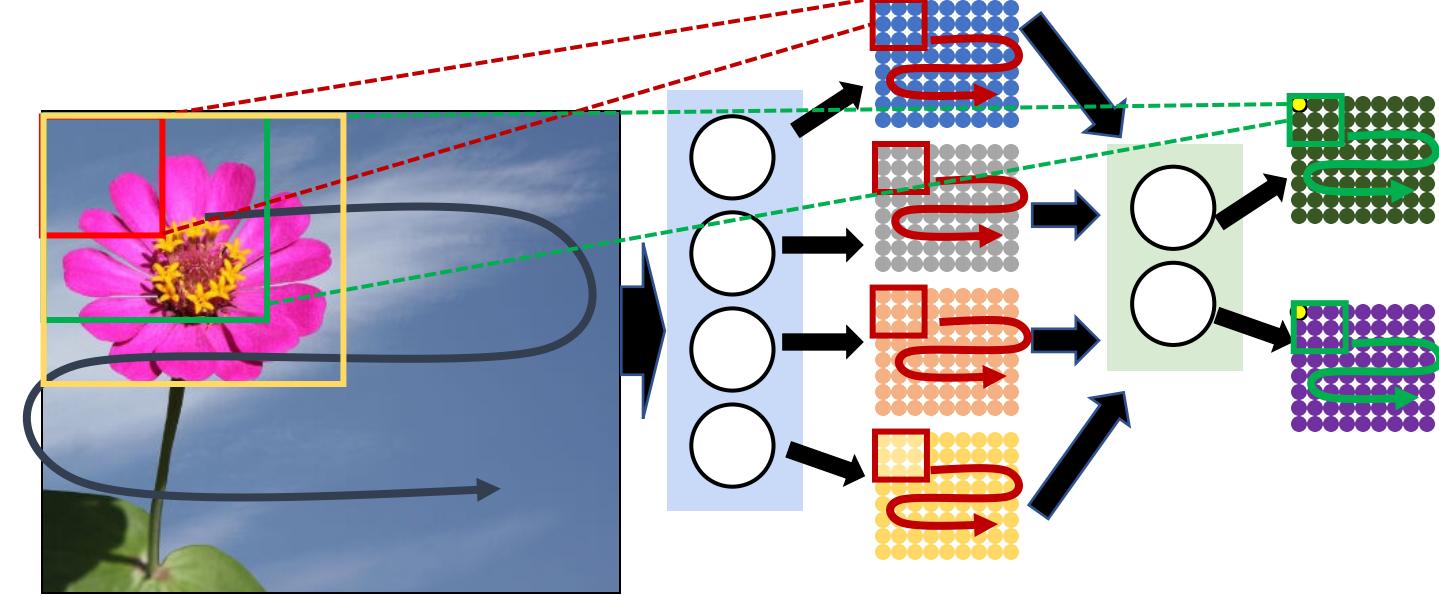
- $\mathcal{O}\left(L_1^2N_1 + L_2^2N_1N_2 + \left(\frac{K}{L_1L_2}\right)^2 N_2N_3 + \dots\right)$
- Here, let $K = 16, L_1 = 4, L_2 = 4, N_1 = 4, N_2 = 2, N_3 = 1$
- Total $64+128+8 = 160$ weights

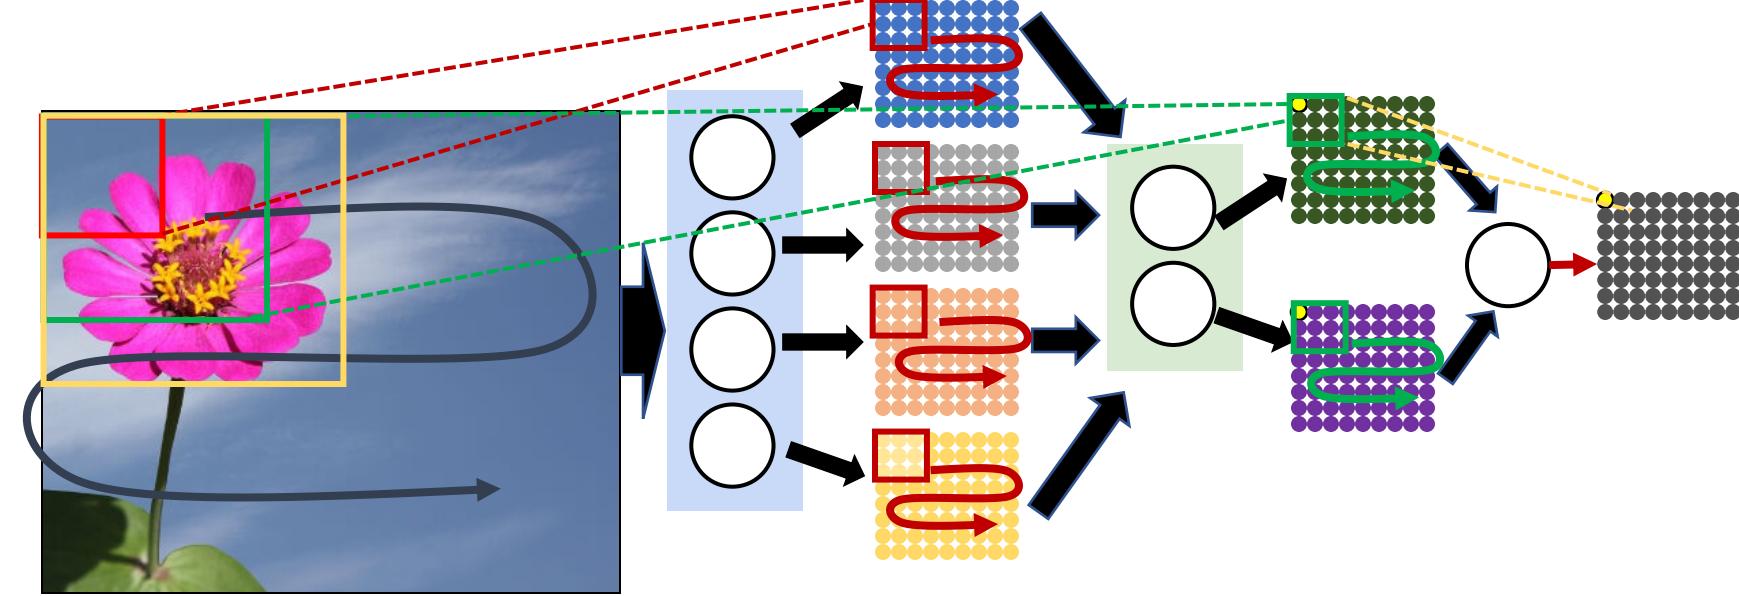


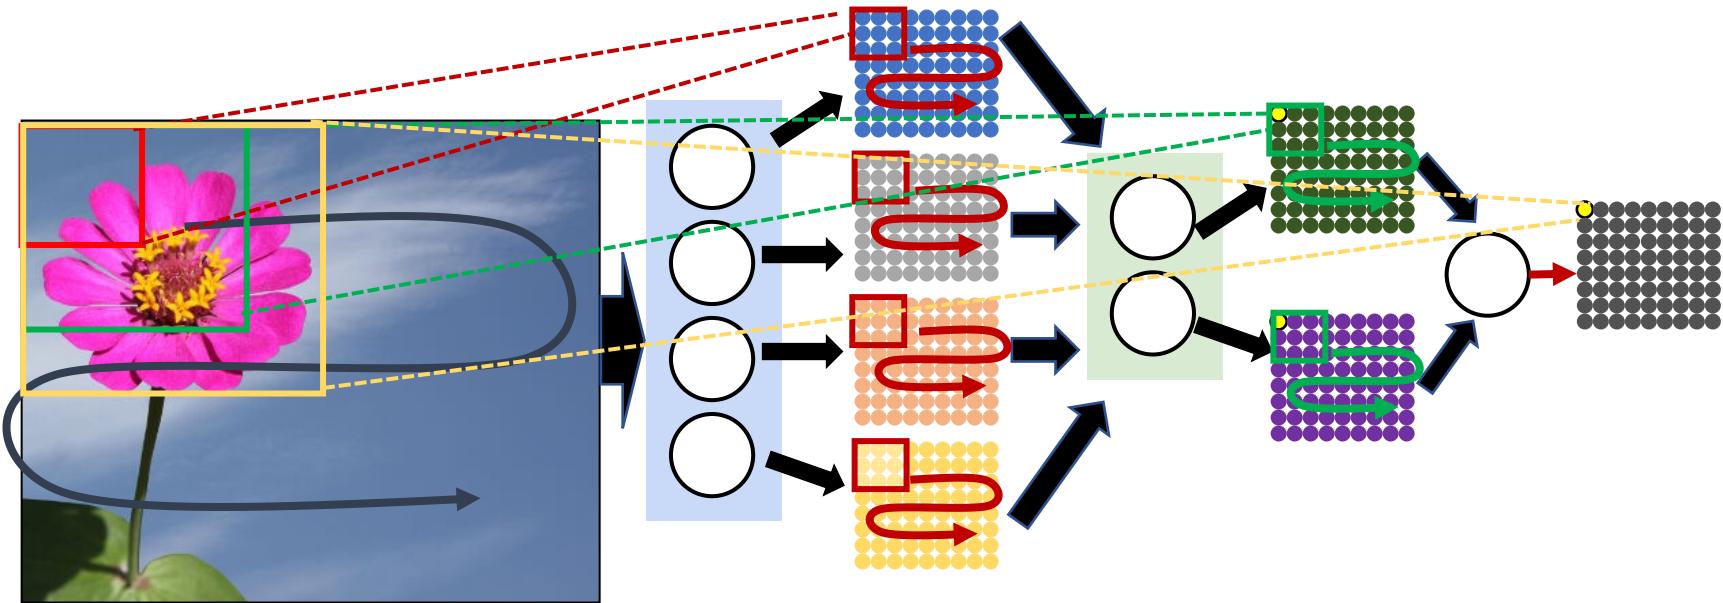




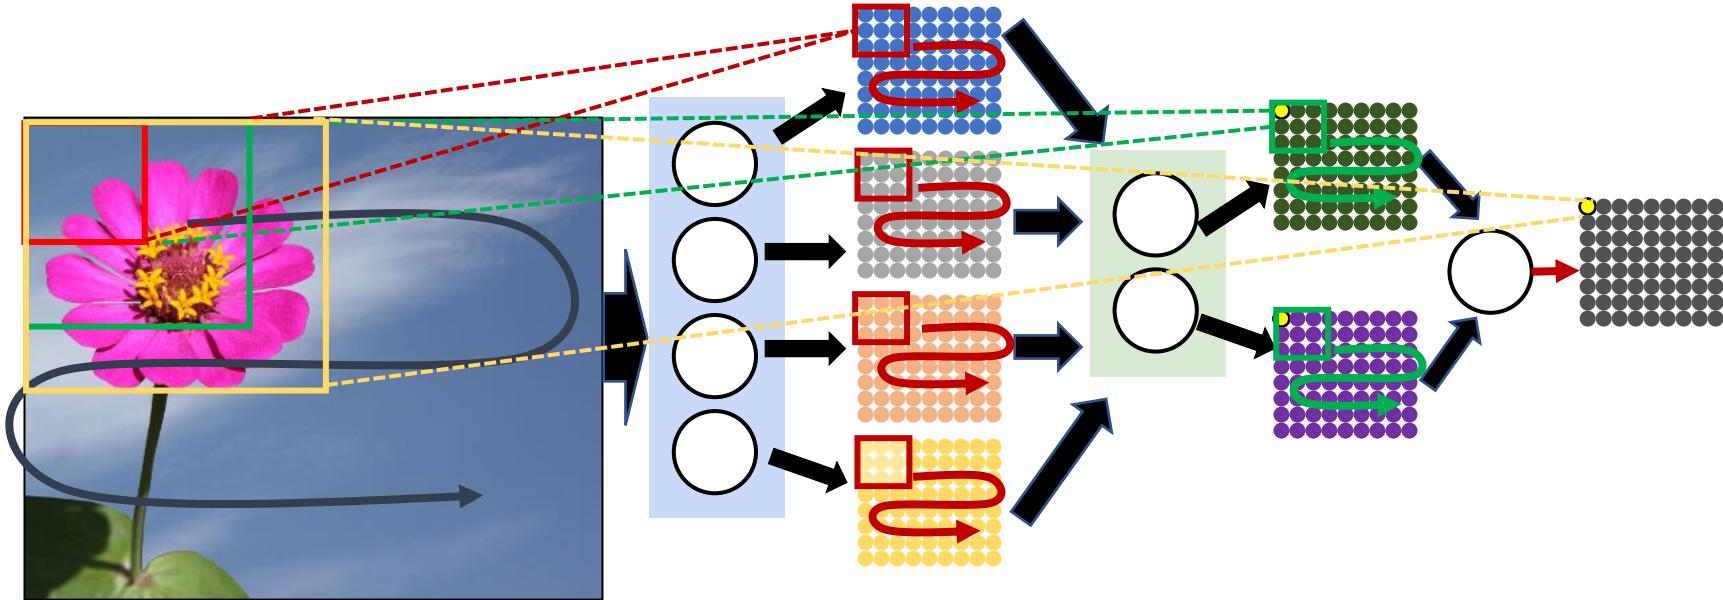








Convolutional Neural Networks (CNNs)

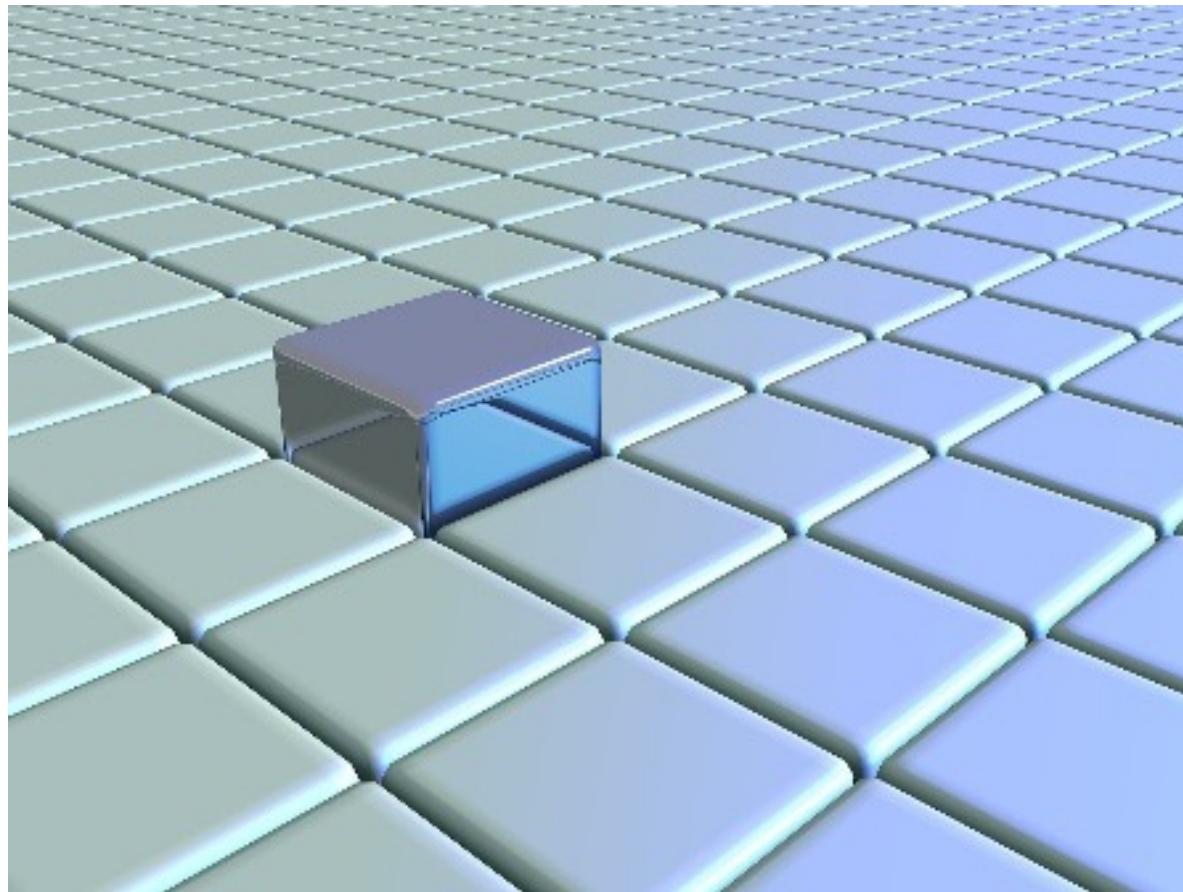


Each of the scanning neurons is generally called a “filter”

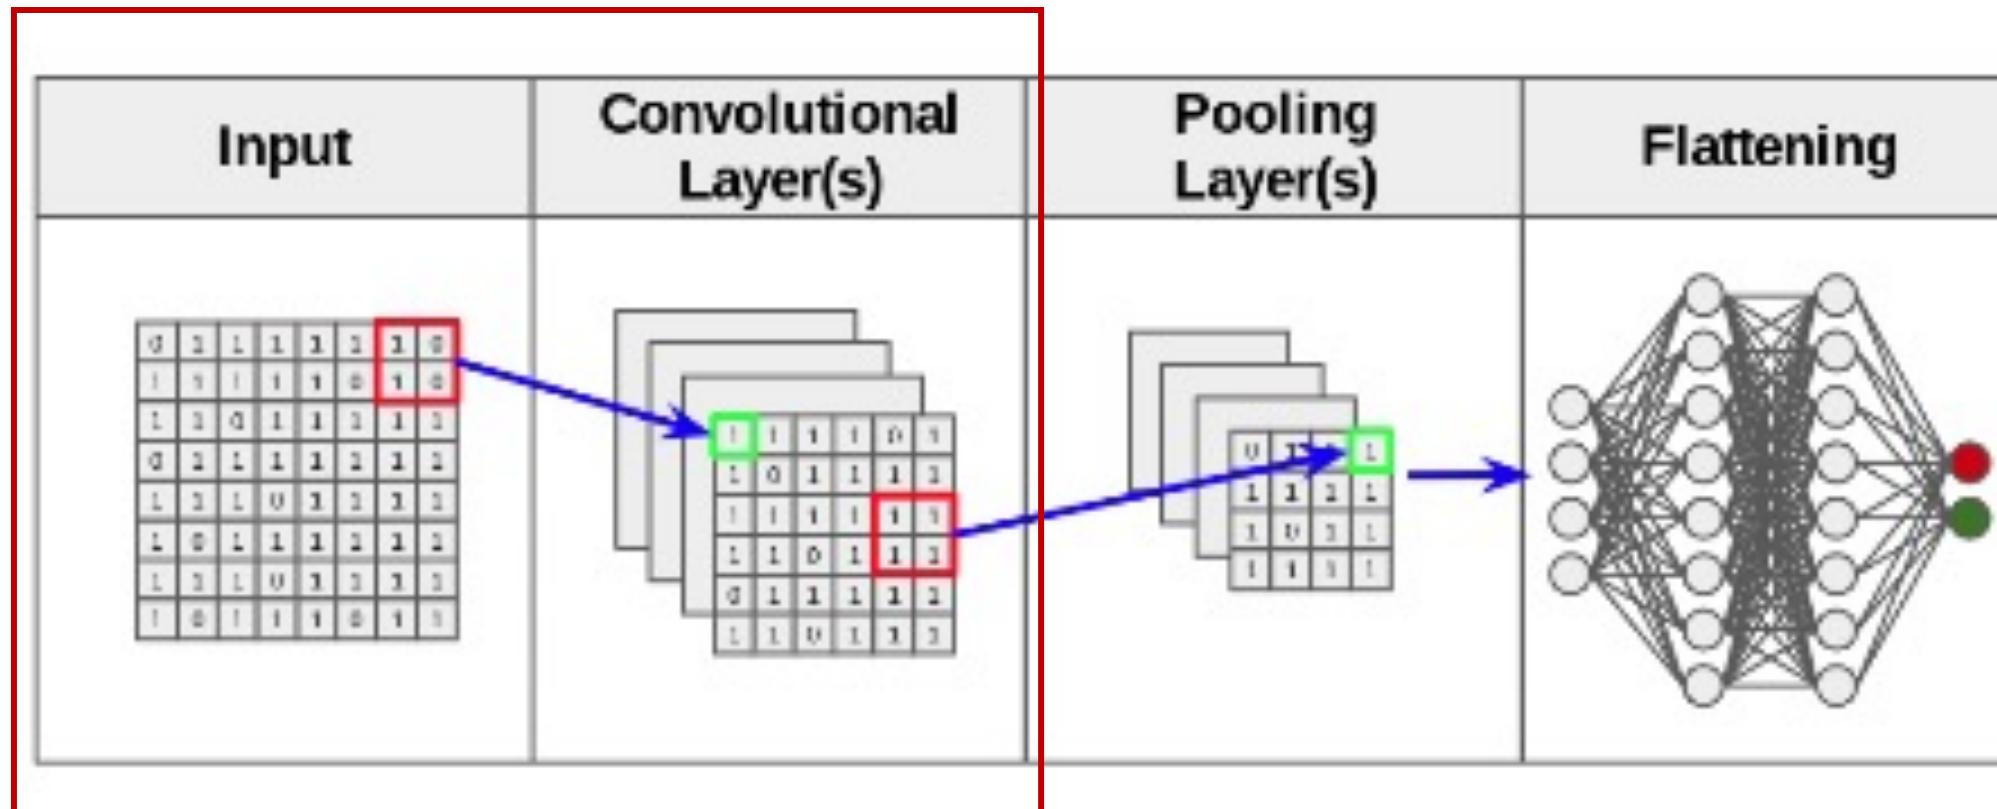
Its really a **correlation filter** as we saw earlier

Each filter scans for a pattern in the map it operates on

For a given neuron, the visual space that affects whether or not that neuron will fire is known as its "**receptive field**."



Convolutional Neural Networks



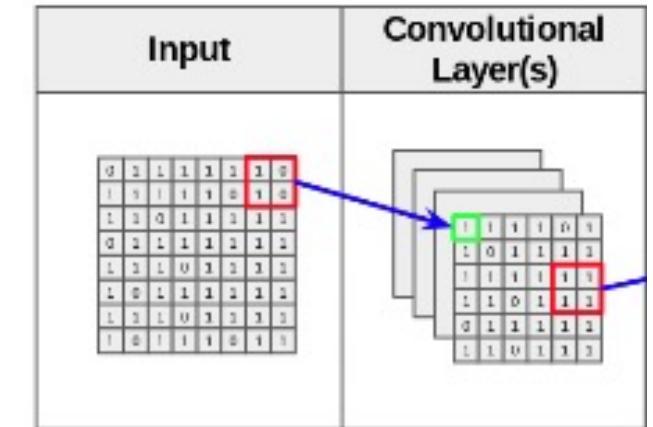
Convolutional Neural Networks: Convolution

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input

1	0	1
0	1	0
1	0	1

Filter



Bias: 0

Convolutional Neural Networks: Convolution

1	0	1
0	1	0
1	0	1

Filter

1 <small>×1</small>	1 <small>×0</small>	1 <small>×1</small>	0	0
0 <small>×0</small>	1 <small>×1</small>	1 <small>×0</small>	1	0
0 <small>×1</small>	0 <small>×0</small>	1 <small>×1</small>	1	1
0	0	1	1	0
0	1	1	0	0

Input Map

4		

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

Input

Filter

1	0	1
0	1	0
1	0	1

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

4	3	

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

4	3	4
2	4	

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

4	3	4
2	4	3
2	3	

1	1	1	0	0
0	1	1	1	0
0	0	1	1	1
0	0	1	1	0
0	1	1	0	0

4	3	4
2	4	3
2	3	4

4	3	4
2		

4	3	4
2	4	3
2		

Blurring

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

Blurring

50	50	100	100
50	50	100	100
50	50	100	100
50	50	100	100
50	50	100	100

Image

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

Kernel
Mask
Filter

New Image

50	50	100	100
50	50	100	100
50	50	100	100
50	50	100	100
50	50	100	100

Image

1/9	1/9	1/9
1/9	1/9	1/9
1/9	1/9	1/9

Kernel
Mask
Filter

New Image

50	50	100	100
50	50	100	100
50	50	100	100
50	50	100	100
50	50	100	100

Image

New Image

$$1/9 \cdot 50 + 1/9 \cdot 50 + 1/9 \cdot 100 \\ + 1/9 \cdot 50 + 1/9 \cdot 50 + 1/9 \cdot 100 \\ + 1/9 \cdot 50 + 1/9 \cdot 50 + 1/9 \cdot 100$$

66.6

$$1/9 \cdot 50 + 1/9 \cdot 50 \\ + 1/9 \cdot 50 + 1/9 \cdot 50$$

55.5

66.6

Edge Detection

$$G_x = \begin{array}{|c|c|c|} \hline -1 & 0 & 1 \\ \hline -2 & 0 & 2 \\ \hline -1 & 0 & 1 \\ \hline \end{array}$$

Measures changes in x-direction
(i.e. detects vertical lines)

$$G_y = \begin{array}{|c|c|c|} \hline -1 & -2 & -1 \\ \hline 0 & 0 & 0 \\ \hline 1 & 2 & 1 \\ \hline \end{array}$$

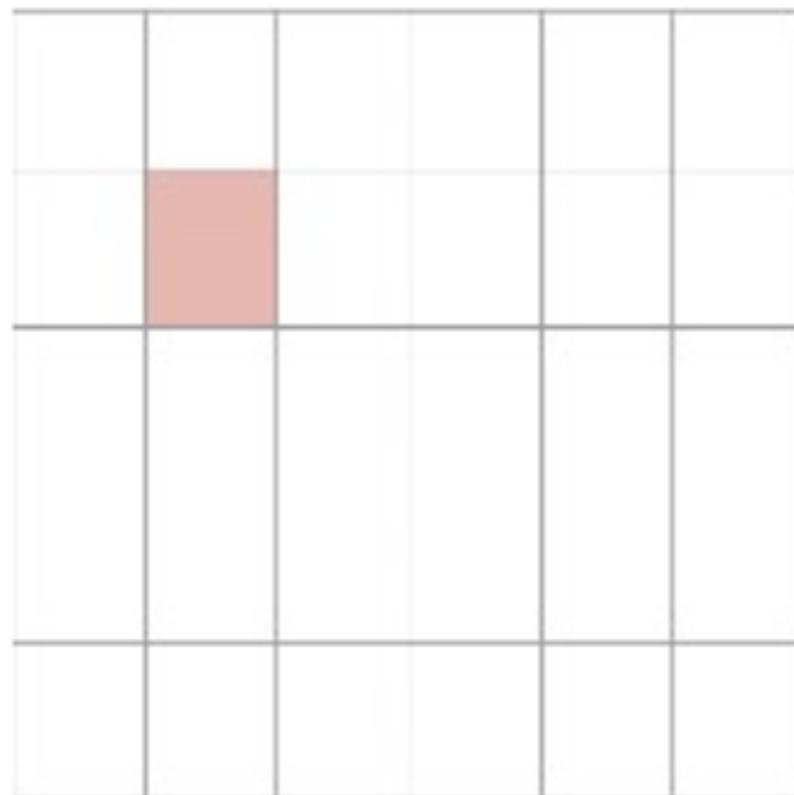
Measures changes in y-direction
(i.e. detects horizontal lines)

Edge Detection

50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100

$$G_x =$$

-1	0	1
-2	0	2
-1	0	1



50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

		0			

50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

		0	200 /9		

50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100

$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

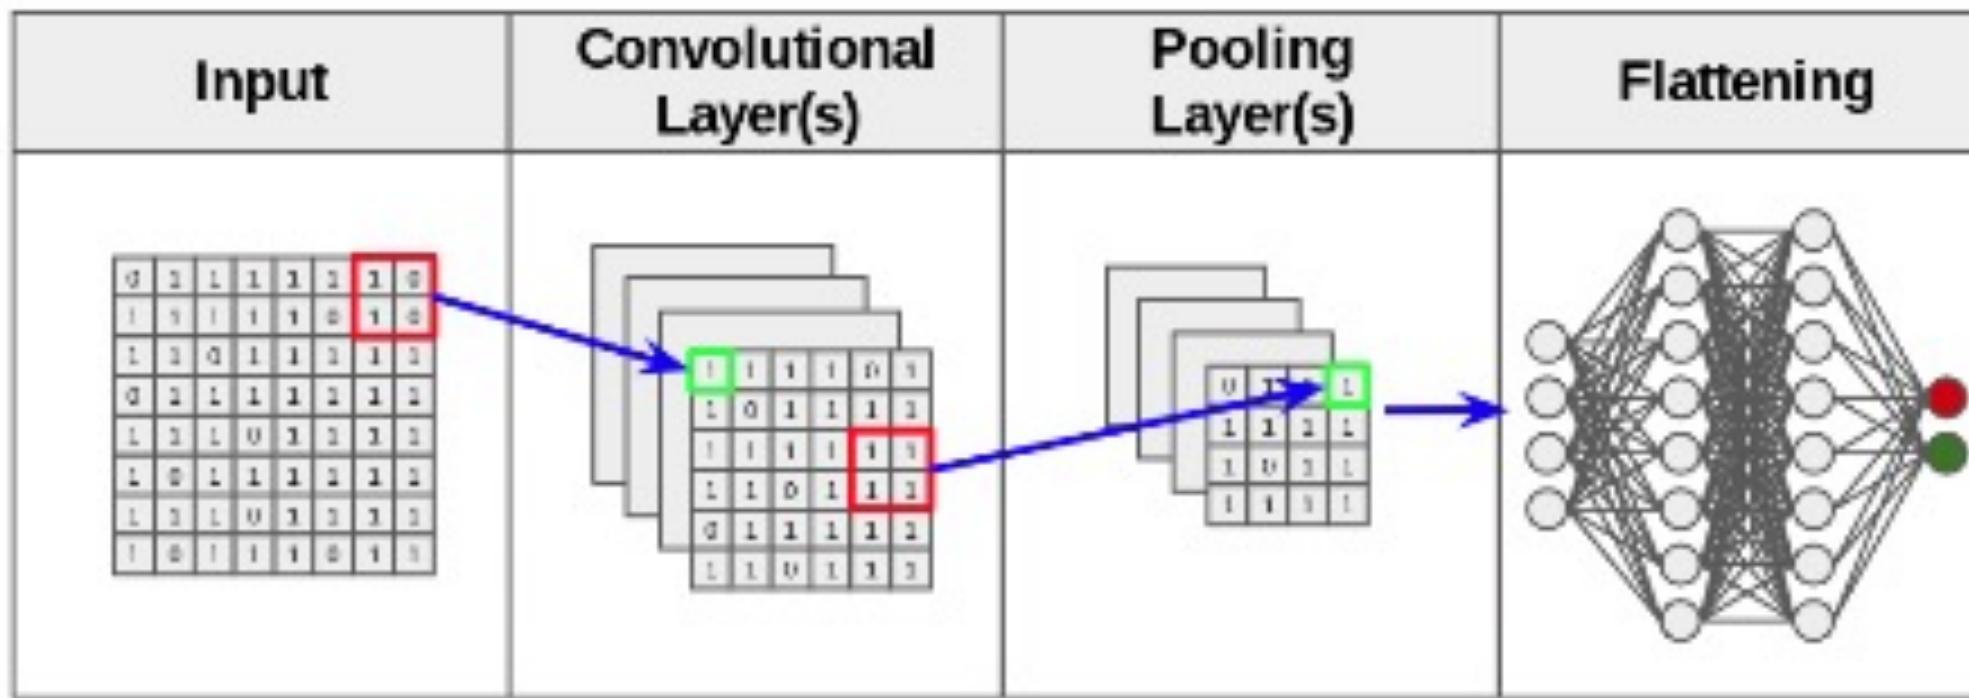
		0	200 /9	300 /9	

50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100
50	50	50	100	100	100

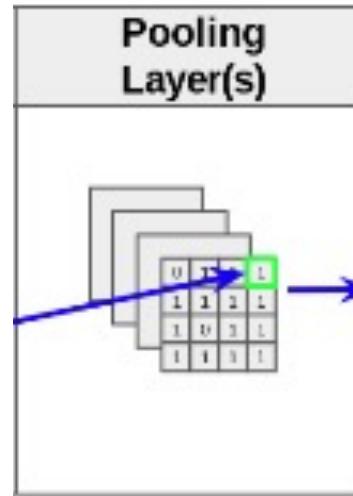
$$G_x = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

		0	200 /9	300 /9	0

Convolutional Neural Networks

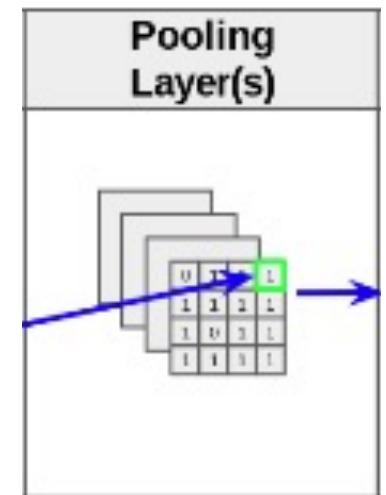


Convolutional Neural Networks: Pooling



Pooling is a type of downsampling that often occurs after convolution. The goal is, without losing much information, to reduce the size of the training data before it goes into the fully connected network.

Convolutional Neural Networks: Pooling

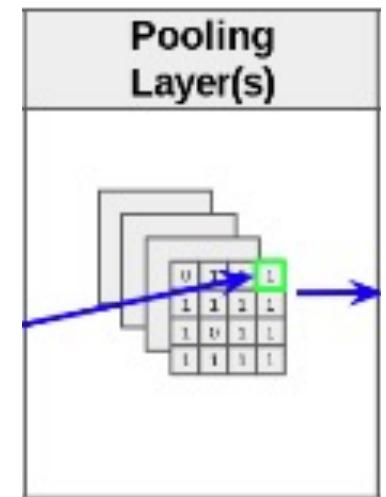


1. Pick a window size
2. Pick a strike
3. Move your window across each of the filtered image
4. Take maximum/minimum, average value in each window

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

Convolutional Neural Networks: Pooling

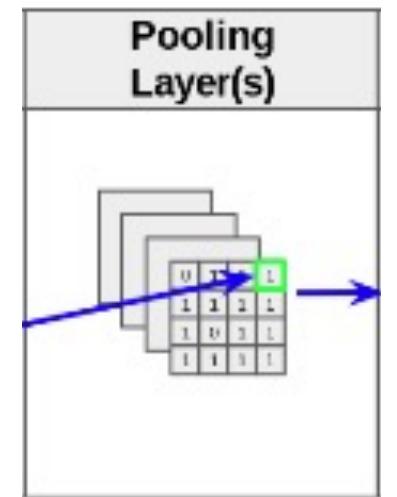


1. Pick a window size
2. Pick a strike
3. Move your window across each of the filtered image
4. Take maximum/minimum, average value in each window

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

Convolutional Neural Networks: Pooling



1. Pick a window size
2. Pick a strike
3. Move your window across each of the filtered image
4. Take maximum/minimum, average value in each window

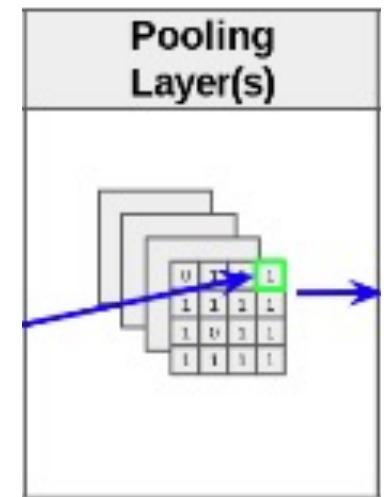
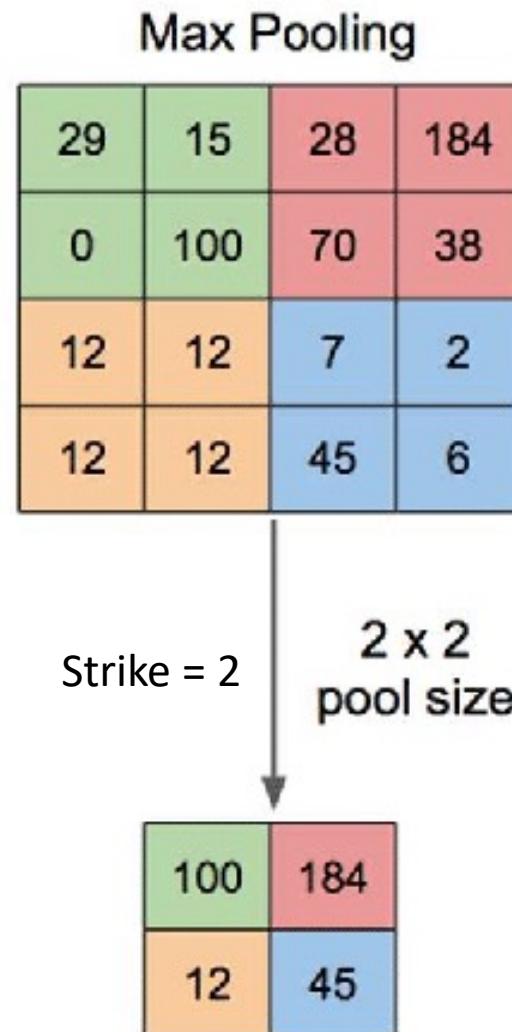
29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

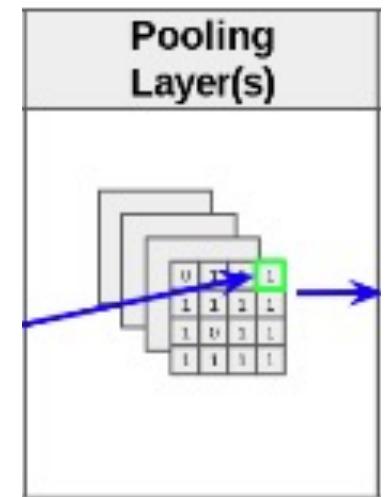
29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

29	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

Convolutional Neural Networks: Pooling



Convolutional Neural Networks: Pooling



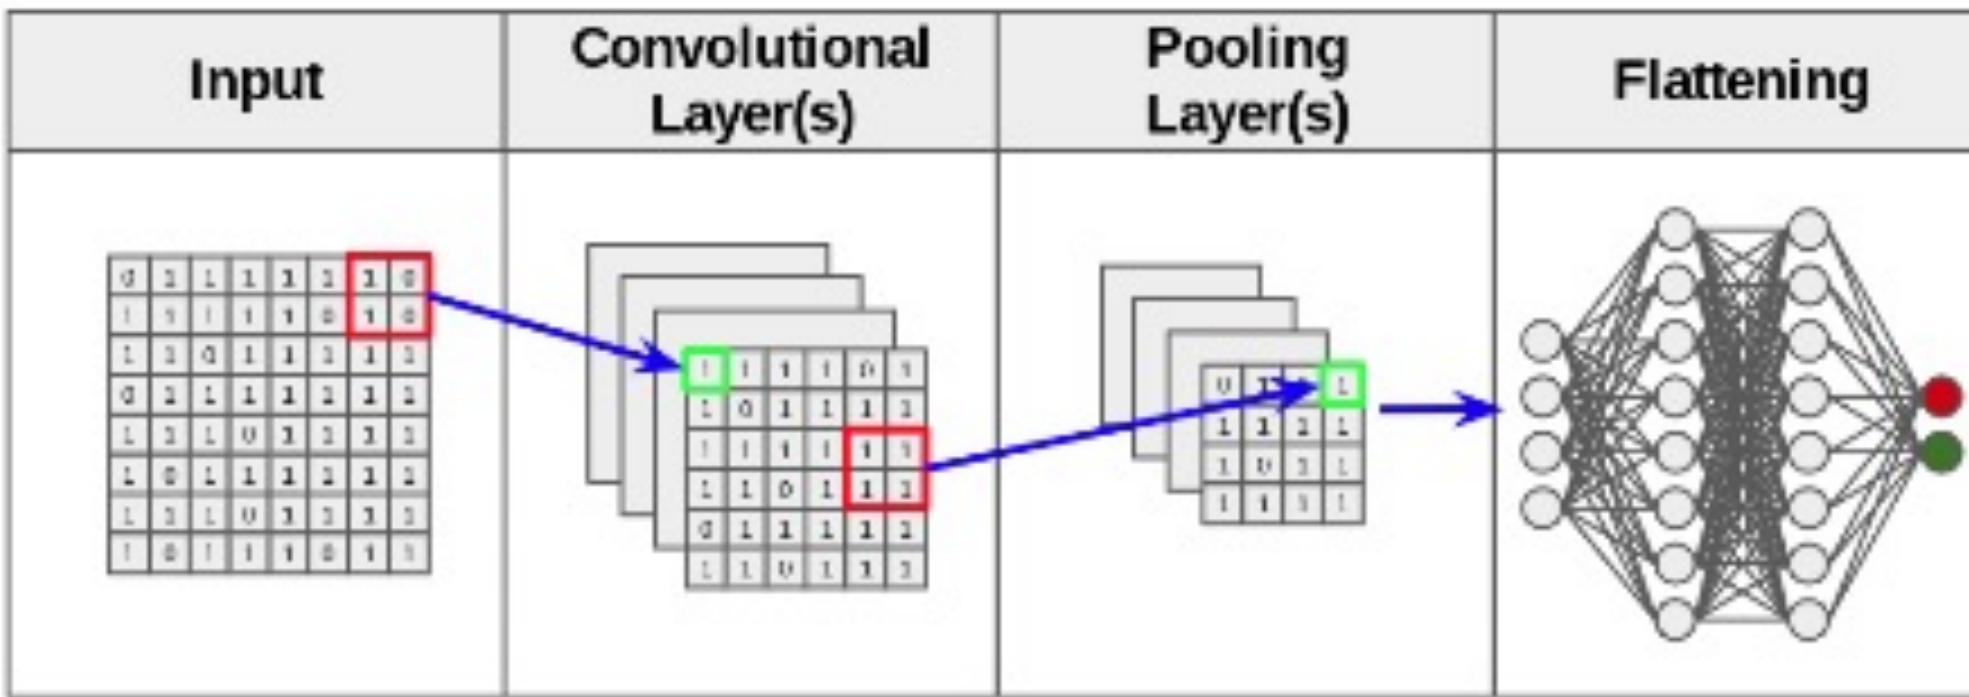
Average Pooling

31	15	28	184
0	100	70	38
12	12	7	2
12	12	45	6

Strike = 2
 2×2
pool size

36	80
12	15

Hyperparameters



How many of each layers and in what orders

Hyperparameters

- Convolution Layer
 - Number of filters
 - Size of filters
- Pooling Layers
 - Window size
 - Strike
- Fully Connected Layers
 - Number of nodes

Your Turn!

- **Exercise 1: Manual Filtering**

```
flower = load_sample_image('flower.jpg')
X = tf.convert_to_tensor([flower], dtype=tf.float32)
convolution = tf.nn.conv2d(X, filters, stride, padding="SAME")
```

Your Turn!

- **Exercise 2: ImageDataGenerator**
- Load data (similar Building a CNN)
- Build a model

```
model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, 3, padding='same', activation='relu',
        input_shape=(160, 160, 3),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(32, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(32, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(32, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

Your Turn!

- **Exercise 2: ImageDataGenerator**
- Load data (similar Building a CNN)
- Build a model

```
model = tf.keras.Sequential([
    tf.keras.layers.Conv2D(32, 3, padding='same', activation='relu',
        input_shape=(28, 28, 3)),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(64, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Conv2D(128, 3, padding='same', activation='relu'),
    tf.keras.layers.MaxPooling2D(),
    tf.keras.layers.Flatten(),
    tf.keras.layers.Dense(128, activation='relu'),
    tf.keras.layers.Dense(10, activation='softmax')
])
```

```
model.compile(optimizer='adam',
              loss='categorical_crossentropy',
              metrics=['accuracy'])
```

Your Turn!

- **Exercise 2: ImageDataGenerator**
- Load data (similar Building a CNN)
- Build a model
- Create an `ImageDataGenerator` to augment the data and then train

```
from tensorflow.keras.preprocessing.image import ImageDataGenerator
```

```
train_data_gen = ImageDataGenerator().flow_from_directory(  
    batch_size = 128,  
    directory=training_dir,  
    target_size=(100, 100))
```

```
model.fit
```

Your Turn!

- **Exercise 2: ImageDataGenerator**
- Load data (similar Building a CNN)
- Build a model
- Create an `ImageDataGenerator` to augment the data and then train
- Once training is complete, we can plot accuracy across time.

```
import matplotlib.pyplot as plt

plt.plot(list(range(len(history.history['accuracy']))),
         history.history['accuracy'])
plt.show()
```

Your Turn!

- **Exercise 2: ImageDataGenerator**
- Load data (similar Building a CNN)
- Build a model
- Create an `ImageDataGenerator` to augment the data and then train
- Once training is complete, we can plot accuracy and loss across time.

```
import matplotlib.pyplot as plt

plt.plot(list(range(len(history.history['accuracy']))),
         history.history['accuracy'])
plt.show()
```

Your Turn!

- **Exercise 2: ImageDataGenerator**
- Load data (similar Building a CNN)
- Build a model
- Create an `ImageDataGenerator` to augment the data and then train
- Once training is complete, we can plot accuracy and loss across time.
- Finally, we want to make predictions using our test dataset.

```
test_image_iterator = tf.keras.preprocessing.image.DirectoryIterator(  
    target_size=(100, 100),  
    directory=test_dir,  
    shuffle=False,  
    image_data_generator=None)  
  
predictions = model.predict(test_image_iterator)
```

Your Turn!

- **Exercise 2: ImageDataGenerator**
- Load data (similar Building a CNN)
- Build a model
- Create an `ImageDataGenerator` to augment the data and then train
- Once training is complete, we can plot accuracy and loss across time.
- Finally, we want to make predictions using our test dataset.
- Print the F1 score

```
f1_score(actual_classes, predicted_classes, average='micro')
```